

C · L · A · I · R · E

# Claude Learns and Improves Iteratively from Real Engagement

An internal AI optimization system built to make Claude measurably better at your specific work — and to understand how.

# Most people use AI like a search engine. That's not its ceiling.

## WHAT MOST PEOPLE EXPERIENCE

Ask a question, get an answer  
Draft an email, clean it up  
Summarize a document  
Generic output — same for everyone  
No memory of past sessions  
Optimized for average, not for you

## WHAT'S ACTUALLY POSSIBLE

- AI that knows your role and context
- Output calibrated to expert-level standards
- Preferences that persist across sessions
- Behavior shaped by how you actually work
- A system that gets better over time
- Competitive advantage, not just convenience

# CLAIRE is not a product. It's a practice — with automation built around it.

**CLAIRE** is a weekly automated pipeline that reads real-world discussions about Claude, identifies patterns relevant to a specific user's work, and surfaces configuration improvements for human review. The human decides what gets applied. Claude gets measurably better. The loop closes.



## OBSERVE

Reads thousands of real user discussions — Reddit, Hacker News, developer forums



## TRIAGE

AI filters signal from noise: real friction vs. one-off complaint



## IMPROVE

Surfaces concrete configuration candidates for human review — evidence-backed

Runs in under 2 minutes · ~\$0.30/week in AI compute · Requires one human decision per candidate

## Why build this at all?

*"AI out of the box is optimized for the average user. Average is the floor, not the ceiling. The goal was to find out how far the ceiling actually goes."*



### **Productivity has a compounding problem**

A 15% efficiency gain from a well-configured Claude compounds differently than a 3% gain from a default one. The configuration IS the investment.



### **Learning from others at scale**

Thousands of people publish what's working and not working with Claude every week. CLAIRE reads them so you don't have to.



### **Institutional risk is real**


Depending on one AI vendor with no documented understanding of how it works is a governance gap. CLAIRE makes the configuration layer explicit.



### **Human judgment, AI leverage**

The system surfaces candidates. Humans make decisions. This is the only safe way to run configuration drift — with a paper trail and a veto.

## Five stages. One loop. Runs weekly.



*FEEDBACK LOOP — Applied changes generate new sessions → new signal → next week's digest*

**~600**

lines of Python

**<2 min**

runtime per cycle

**\$0.30**

per weekly run

## What the system actually produces.

### WEEKLY DIGEST — CANDIDATE #3 OF 7 — TRACK A: CONFIGURATION EDIT

**SIGNAL TYPE:**

Response Format — Citation / Verification Behavior

**EVIDENCE:**

14 posts across Reddit (r/ClaudeAI) and HN flagging Claude confidently citing unverifiable figures without qualification. Pattern consistent across 3 consecutive weeks.

**CANDIDATE EDIT:**

Add to System Prompt: "For any statistic, regulatory reference, or named source — if confidence is below certainty, append [verify] inline. Never assert uncertain facts as confirmed."

**CONFIDENCE:**

HIGH (0.87) — Corroborated by 14 posts, signal type matches known friction area, no conflicting evidence

HUMAN DECISION REQUIRED →

APPLY

DEFER

REJECT

*All decisions require a human-written hypothesis. Every change is logged with author + rationale.*

## CLAIRE vs. CLAIRE-A: The most important distinction in the system.

### CLAIRE (Original)

- Agentic pipeline — structured automation
- Ingests, triages, synthesizes on a schedule
- Outputs candidates for human review
- Human writes the hypothesis
- Human decides: Apply · Defer · Reject
- No autonomous action on live configuration
- Audit trail on every decision

### CLAIRE-A (Shadow Layer)

- Autonomous agent layer — but shadow-only
- Reads candidate payload from CLAIRE output
- Opus model reasons, writes its own hypothesis
- Sonnet model scores the decision quality
- Builds a reliability ledger over time
- Writes NOTHING to live configuration
- Every output is auditable, reversible, logged

*CLAIRE-A is how we study what autonomous AI decision-making looks like — before we ever let it act.*

**"Agentic" means AI does the work.**

**"Autonomous" means AI makes the decisions.**

### AGENTIC PIPELINE (CLAIRE)

AI

Reads 500+ posts

AI

Scores relevance

AI

Writes candidate summary

HUMAN

Decides to apply or defer

HUMAN

Writes hypothesis / rationale

HUMAN

Logs and applies the change

**Human remains decision-maker.**  
**AI does the research and drafting.**

### AUTONOMOUS AGENT (CLAIRE-A, SHADOW)

AI

Reads candidate payload

AI

Reasons about evidence

AI

Writes its own hypothesis

AI!

Decides: Apply / Defer / Reject

AI

Scores quality of own decision

LOG

Logs — but acts on nothing

**AI makes decisions — but only in shadow.**  
**Nothing touches live config without human approval.**

# Keeping humans in the loop isn't a limitation. It's the point.

*"Automated configuration drift with no accountability is how you end up with a Claude that's been optimized by the internet into something unrecognizable."*



## Governance

Every change has an author, a rationale, and a date. When something goes wrong, you know exactly why — and who approved it.



## Judgment

AI can surface evidence. It cannot weigh organizational context, fiduciary obligation, or whether a Reddit-validated change is right for a regulated firm.



## Learning

The hypothesis written before applying a change becomes the baseline for evaluating it. That loop only works if a human is authoring the intent.

# Five things CLAIRE proved about using AI seriously at work.

01

## Configuration is the competitive moat

The model is the same for everyone. How you configure it is not. That gap compounds over time.

02

## AI learns from other AI users — if you let it

Thousands of people document what works and fails every week. The insight is available. The question is whether you're reading it.

03

## Automation does not equal autonomy

A pipeline that does the research work is not the same as one that makes the decisions. The distinction matters for governance and trust.

04

## The hypothesis is the product

The most durable output isn't the config change — it's the written, testable rationale attached to it.

05

Shadow before live — always — CLAIRE-A runs autonomous decisions against every candidate and acts on none. You don't deploy autonomous AI until you've watched it decide for months.

# CLAIRE's name is recursive on purpose.



## Community Signal

Reddit · HN · dev.to  
discussions surface  
patterns



## AI Triage & Synthesis

Signal scored, candidates  
generated per week



## Human Decision

Review · Hypothesize  
Apply or defer each  
change



## Better Sessions

Claude improves · better  
output → better signal

← LOOP: Better sessions generate better signal → next week's candidates improve → Claude improves further →

*This is why the name is recursive — the system improves itself. Better configuration → better output → better signal → better configuration.*

# Build 7 is live. Three open threads.

1

## First CLAIRE-A Eval Cycle

Timeline: 14–21 days

CLAIRE-A has been making shadow decisions since Build 5. In two to three weeks, enough time will have passed to score whether its autonomous decisions aligned with human ones. First real data on autonomous AI decision quality.

2

## Packaging CLAIRE for Other Users

Timeline: Next design session

Stripping CLAIRE to a Claude Project system prompt — no Python, no local infrastructure. A stranger creates a Project, pastes the prompt, types 'Run CLAIRE', and gets a weekly digest. Zero code installation.

3

## Build 8 — Expanding Signal Sources

Timeline: After packaging

Adding X/Twitter and Substack RSS feeds to the ingest layer. Targets the gap in regulated-industry AI use discussion. More signal diversity means better candidates.

## The Thesis

# AI is not a tool you buy. It's a capability you build.

**CLAIRE** is proof that a focused practitioner — with \$0.30 a week in AI compute and a commitment to one weekly review — can build a version of Claude measurably better calibrated to their specific domain, workflow, and standards than the version every other subscriber gets.

*That's the value. The pipeline is just the mechanism.*