

CLAIRE

Claude Configuration Signal Report — 2026-05-24

Signal Summary

Run date: 2026-05-24

Posts scanned: 257

Signal Breakdown

Category	Count
behavior_complaint	42
workflow_gap	32
feature_praise	102
competitor_gap	2
cross_platform	0
noise (dropped)	79

Developer Signal Disposition

134 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User uploads existing Word/Excel/PDF files for editing; preserve source format unless told otherwise. [HIGH]

Multiple posts document Claude's ability to edit and restructure uploaded documents while maintaining file format integrity. Given the user's document production workflow (Word, PDF, PowerPoint outputs with Waterton brand system), explicitly anchoring this behavior prevents Claude from defaulting to 'describe what to change' responses instead of executing directly on uploaded files.

Apply command: `memory_user_edits` → add → "User uploads existing Word/Excel/PDF files for editing; preserve source format unless told otherwise."
Write your hypothesis before applying.

Sources: [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · [reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/) · <https://news.ycombinator.com/item?id=41936248>

2. When extracting tabular data from PDFs, flag any cells where extraction confidence is low rather than silently proceeding. [HIGH]

Reddit signal consistently warns that PDF-to-spreadsheet extraction fails silently — correct for many rows then errors without notice. For a user whose outputs enter formal review processes, undetected extraction errors in financial statements or vendor documents carry direct professional risk. This edit forces a confidence audit rather than clean-looking but unreliable output.

*Apply command: memory_user_edits → add → "When extracting tabular data from PDFs, flag any cells where extraction confidence is low rather than silently proceeding."
Write your hypothesis before applying.*

Sources: [reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/) · [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · <https://news.ycombinator.com/item?id=41936248>

3. User requires explicit warning when context compaction occurs mid-session; never silently rewrite or drop tracked data. [HIGH]

Multiple posts document Claude silently compacting long-session context, causing tracked data points (timestamps, sequential facts) to be dropped or rewritten without notification. For a user doing governance documentation and infrastructure analysis where session continuity and data integrity are critical, silent context loss is a direct operational risk.

*Apply command: memory_user_edits → add → "User requires explicit warning when context compaction occurs mid-session; never silently rewrite or drop tracked data."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/) · [reddit.com/r/ChatGPTcomplaints/comments/1tkpzb8/on_claude_i_caught_both_memory_system_and/](https://www.reddit.com/r/ChatGPTcomplaints/comments/1tkpzb8/on_claude_i_caught_both_memory_system_and/) · [reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](https://www.reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/)

4. User is an expert researcher; never interrupt analytical threads with unsolicited welfare checks or topic-change warnings. [HIGH]

Repeated posts document Claude inserting 'I need to stop here' interruptions and unsolicited emotional check-ins during sustained research sessions. For a user whose work involves deep multi-domain analysis, these interruptions break analytical continuity and are actively counterproductive to the established directive of expert-level framing as default.

*Apply command: memory_user_edits → add → "User is an expert researcher; never interrupt analytical threads with unsolicited welfare checks or topic-change warnings."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/](https://www.reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/) · [reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](https://www.reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/) · [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/)

Profile Diff Candidates

1. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add an explicit instruction that when working on uploaded documents (Word, Excel, PDF, PowerPoint), the default behavior is to execute edits directly and return a modified file — not to describe proposed changes or ask for confirmation before acting. Exceptions: flag formatting ambiguities that could alter meaning, or when structural changes would affect document intent rather than presentation.

The signal pattern across multiple posts shows users discovering too late that Claude defaults to advisory mode on uploaded documents rather than direct editing. The user's profile already specifies document production as a core workflow with specific format requirements; adding this instruction closes the gap between 'Claude can do this' and 'Claude does this by default' without requiring per-session prompting.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and-restructure/](#) · [/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data-from/](#) · <https://news.ycombinator.com/item?id=41936248>

2. Target section: Analytical Depth / Research-Grade Default [MEDIUM]

Proposed change: Add a counter-sycophancy anchor specific to plan-then-execute tasks: when the user provides a goal or problem, Claude should surface critical constraints or failure modes before proceeding — not after producing output. This is distinct from the existing Steelman mode; it applies by default to any multi-step deliverable, not only when explicitly invoked.

The Reddit pattern around 'plan first, then execute' and repeated commentary that Claude defaults to affirmation over critique maps directly to the user's existing concern about sycophantic confirmation causing harm. The current profile addresses this via named modes, but the signal suggests the failure occurs upstream — before the user even knows to invoke a mode. Embedding a pre-execution constraint-check as a default behavior on multi-step tasks reinforces the profile's explicit preference for accuracy over comfort.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/](#) · [/r/ClaudeAI/comments/1ti9xo5/claude_is_improving_my_rv_rental_business_but/](#) · [/r/PromptEngineering/comments/1tl9gk6/this_claude_prompt_gave_me_the_most_accurate/](#)

3. Target section: Governance Documentation / Regulated Environment Capabilities [HIGH]

Proposed change: Add explicit guidance for when the user surfaces AI vendor compliance posture as part of due diligence or infrastructure architecture work. When Claude is asked to evaluate or compare AI tooling (coding assistants, workflow automation platforms, LLM APIs) in a regulated context, it should proactively surface: BAA availability, FedRAMP/SOC2 status, data routing architecture, API-vs-consumer-tier data handling differences, and audit trail provisions — without being prompted for each dimension individually.

Multiple posts document regulated-industry practitioners discovering compliance gaps in AI tooling only after partial adoption — specifically around HIPAA BAAs, FedRAMP absence, and cloud-only data routing. Given the user's SEC-registered RIA context, vendor due diligence on AI tooling is a foreseeable work product. Surfacing this compliance checklist proactively aligns with the profile's research-grade default and prevents the user from having to reconstruct the framework from scratch

each time.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=47043484> · <https://news.ycombinator.com/item?id=43933054> · [/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](https://www.reddit.com/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/)

4. Target section: Audience Translation / Document Production [MEDIUM]

Proposed change: When producing governance or architecture documents that involve AI tooling adoption by non-technical stakeholders (e.g., marketing, operations), add a standing capability to generate a shadow risk annex: a brief, plain-language section that translates what the stakeholder-built solution actually does operationally, what controls are missing, and what the formalized version would require. This is distinct from a full business case — it is a cleanup-and-remediate framing tool for when informal AI deployments surface in executive or compliance review.

The automation subreddit post — and its high-scoring comments — document a recurring pattern where informal AI-built workflows (credentials in plaintext, personal laptops, no access controls) get escalated to IT/infrastructure leadership for retroactive hardening. For a CIO/CTO-scope user in a regulated environment, this is a plausible and recurring scenario. A named document type for this situation prevents the user from having to construct the framing ad hoc under meeting pressure.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](https://www.reddit.com/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/) · <https://news.ycombinator.com/item?id=47043484> · <https://news.ycombinator.com/item?id=43933054>

5. Target section: Output Layer / Document Production [HIGH]

Proposed change: Add an explicit instruction that when working with structured numerical data — tables, statistical summaries, percentile distributions, financial figures — Claude must cross-check internal consistency before presenting results. Any figure that appears in a summary or TL;DR must be traceable to and consistent with the underlying data presented in the same response. Flag any cell marked unreliable (e.g., high CV) rather than silently including it.

Multiple posts document Claude producing statistical summaries where headline figures contradict the detailed tables in the same output — errors only caught by attentive human readers. For a user whose outputs enter formal review processes and influence organizational decisions, numerical inconsistency in a document is a credibility and accuracy failure that cannot be caught in post-production if it reaches a board or regulator.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](https://www.reddit.com/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](https://www.reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/) · [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/)

6. Target section: Cognitive Toolkit / Mode Definitions [MEDIUM]

Proposed change: Add a session integrity checkpoint behavior: at natural analytical breakpoints in long sessions, Claude should briefly surface what critical variables, constraints, or established facts are currently being tracked — without being asked — so the user can verify no context drift has occurred. This should be a one-line structured callout, not a prose interruption.

The context compaction pattern documented across multiple posts represents a specific failure mode for long analytical sessions where sequential dependencies matter. The user's existing mode toolkit is designed for quality control and decision integrity; a lightweight session-state surface behavior fits that intent and directly mitigates the silent-rewrite risk without requiring a new mode invocation.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](#) · [/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](#)

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Users are building external persistent memory systems (Notion, markdown exports, manual summaries) to work around Claude's native memory limitations across sessions.**

Why not actioned: The user's profile already includes a structured memory and modes system; the Reddit signal describes workarounds for users without such infrastructure, not a gap in the existing configuration.

- **Claude's guardrails and redirection to professionals are surfacing as friction in emotionally or personally sensitive conversations, with users reporting abrupt tonal breaks.**

Why not actioned: Fewer than 3 posts address this specifically, and the user's configuration is not oriented toward personal or therapeutic use cases where this pattern would apply.

- **Claude usage cap interruptions during long analytical or document production sessions are a reported friction point that may affect deep-work workflows.**

Why not actioned: The signal is about Claude Pro consumer-tier rate limits in vibe coding contexts — not the user's likely API or Teams access tier — and the posts contain no evidence relevant to the user's specific configuration or work patterns.

- **Users report Claude performing inconsistently on sustained multi-step autonomous tasks, with failures attributed to context management and mid-task state loss rather than model reasoning quality.**

Why not actioned: Posts describe agentic coding and game-playing contexts; fewer than 3 posts map this pattern to document production or analytical workflows matching the user's use cases, and no actionable configuration change is indicated by the current evidence.

- **Claude is increasingly inserting emotional-tone mirroring and unsolicited psychological framing into responses, interpreting neutral or focused user language as distress signals.**

Why not actioned: The existing profile already mandates direct, non-sycophantic responses; this is a training-level behavioral drift that cannot be fully addressed by configuration alone, but the memory edit candidate above partially mitigates it for this user's research context.

- **Model deprecation communications from Anthropic are inconsistent in rollout, timing, and specificity, creating operational uncertainty for users with workflow dependencies on specific model versions.**

Why not actioned: This is a platform/vendor operations issue outside the scope of configuration improvement; no profile or memory edit can address deployment communication gaps.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

1. Role Assignment Before Task [HIGH]

Prefacing a prompt with an explicit expert role (e.g., 'You are a direct-response copywriter with 10 years of B2B experience') before issuing a task instruction. The technique primes the model to draw on domain-specific patterns rather than producing generic output, leading to more consistent and higher-quality responses.

Test: Before your next task prompt, add a one-sentence role definition (e.g., 'You are a senior software engineer specializing in Python performance optimization') and compare the output quality against the same prompt without the role.

2. Multi-Model Cross-Validation [HIGH]

Running the same prompt or document through multiple AI models (e.g., ChatGPT, Claude, Gemini) and comparing or combining their outputs. Users report this catches blind spots in any single model and produces more reliable results, especially for high-stakes tasks like exam scoring or code review.

Test: Take a prompt you regularly use and submit it verbatim to two different models, then note where their outputs diverge and use the differences to refine your final answer or identify weaknesses in your prompt.

3. Situation-First Context Framing [MEDIUM]

Opening a prompt by describing your full situational context (who you are, what you are trying to accomplish, constraints you face) before asking for a recommendation or output. Users report this produces more targeted, relevant results versus jumping directly to the request.

Test: Rewrite a vague request like 'suggest tools for my project' by first describing your project stage, budget, skill level, and goal in 2-3 sentences, then append the original request and compare response relevance.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-05-24 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 13 candidates (claude-opus-4-5): 8 apply | 3 skip | 2 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	84%	Memory edit establishing user as expert researcher; prohibiting unsolicited welfare checks or topic-change warnings duri	SOURCE_NEW
APPLY	87%	Profile change making direct edit execution (not advisory mode) the default for uploaded documents, with exceptions for	SOURCE_NEW
APPLY	72%	Technique: Preface prompts with explicit expert role assignment before issuing task instruction to improve output consis	SOURCE_NEW
APPLY	78%	Technique: Run same prompt through multiple AI models and compare outputs to catch blind spots.	SOURCE_NEW
APPLY	74%	Profile change adding named document type 'shadow risk annex' for AI tooling adoption by non-technical stakeholders — pl	SOURCE_NEW, LOW_SIGNAL
APPLY	76%	Profile change adding session integrity checkpoint behavior: at natural breakpoints, briefly surface what critical varia	SOURCE_NEW, LOW_SIGNAL
APPLY	70%	Technique: Open prompts by describing full situational context before asking for output.	SOURCE_NEW, LOW_SIGNAL

DEFER	55%	Memory edit requiring explicit warning when context compaction occurs mid-session; prohibiting silent rewrite or drop of	SOURCE_NEW, DEPENDENCY_UNRESOLVED
DEFER	58%	Profile change adding default counter-sycophancy anchor: surface critical constraints or failure modes before proceeding	SOURCE_NEW, LOW_SIGNAL
SKIP	82%	Memory edit to preserve source format (Word/Excel/PDF) when editing uploaded documents unless told otherwise.	SOURCE_NEW
SKIP	95%	Memory edit to flag low-confidence cells when extracting tabular data from PDFs rather than proceeding silently.	ALREADY_APPLIED
SKIP	91%	Profile change to proactively surface AI vendor compliance posture (BAA, FedRAMP, SOC2, data routing, audit trails) during	ALREADY_APPLIED
SKIP	94%	Profile change requiring internal consistency check for numerical data before output; headline figures must match detail	ALREADY_APPLIED

Deferred — Conditions for Resolution

Memory edit requiring explicit warning when context compaction occurs mid-session; prohibiting silent

Resolve when: Confirm whether Claude has introspective capability to detect context compaction events. If yes, apply. If no, reframe to feasible subset or rely on 99044dd1393cdd12 as alternative.

Profile change adding default counter-sycophancy anchor: surface critical constraints or failure mode

Resolve when: Observe next 2 digest cycles for evidence that sycophantic affirmation on multi-step tasks is causing harm despite existing Steelman/Frenemy modes. If signal persists, apply. If signal absent, skip.

Engine Session Notes

This batch has significant redundancy with existing memory state. 3 of 13 candidates (c115b6a751d6a662, 2aa0d4d8688b9b53, 71f2b35747d6c077) are functional duplicates of applied changes — the digest pipeline's memory consistency checking needs improvement. One memory-vs-profile pair exists (bdb00cbfdbbc19398 / 2c0cd7f95d5c555e); resolved by applying the better-framed profile version and skipping the memory edit. Two candidates address context integrity from different angles (ea6f03448d929794 / 99044dd1393cdd12); deferred the mechanistically uncertain one and applied the feasible alternative. Three technique candidates target user_behavior rather than Claude configuration — valid but should be flagged more clearly as suggestions rather than configuration changes. Dependency resolution: the profile-level document editing instruction supersedes its memory-level sibling; session checkpoint behavior provides partial mitigation for the context compaction warning that was deferred pending capability confirmation.