

CLAIRE

Claude Configuration Signal Report — 2026-05-24

Signal Summary

Run date: 2026-05-24

Posts scanned: 257

Signal Breakdown

Category	Count
behavior_complaint	45
workflow_gap	27
feature_praise	93
competitor_gap	2
cross_platform	1
noise (dropped)	89

Developer Signal Disposition

129 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User uploads PDFs and structured documents for data extraction; always verify extraction fidelity before downstream use. [HIGH]

Multiple posts document Claude's PDF-to-structured-data capability alongside a consistent warning pattern: extraction is accurate most of the time but produces silent errors on edge cases. For a user whose outputs enter formal review processes, undetected extraction errors in financial statements or contracts are high-stakes. This edit surfaces a verification reminder at the point of use.

Apply command: `memory_user_edits` → add → "User uploads PDFs and structured documents for data extraction; always verify extraction fidelity before downstream use."
Write your hypothesis before applying.

Sources: reddit.com/r/PromptEngineering/comments/1tgfwqi/i_didnt_realise_claude_could_extract_data_from/ · <https://news.ycombinator.com/item?id=41343826> · reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/

2. User works with existing Word and Excel files; Claude should offer in-place restructure/edit, not rebuild from scratch. [HIGH]

Posts reveal a widespread discovery gap: users treat Claude as a document generator rather than a document editor, defaulting to rebuilding when upload-and-restructure is available. Given the user's Waterton document production system and formal output requirements, defaulting to edit-in-place preserves formatting integrity and reduces rework cycles.

*Apply command: memory_user_edits → add → "User works with existing Word and Excel files; Claude should offer in-place restructure/edit, not rebuild from scratch."
Write your hypothesis before applying.*

Sources: reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/ · <https://news.ycombinator.com/item?id=41936248> · reddit.com/r/PromptEngineering/comments/1tgfwqi/i_didnt_realise_claude_could_extract_data_from/

3. User requires Claude to flag data accuracy concerns before presenting summarized statistics. [HIGH]

Multiple posts document Claude presenting quantitatively inconsistent summaries of structured data without flagging contradictions — notably the ONS income data post where median figures were internally inconsistent. In a regulated RIA environment where financial and statistical outputs enter formal review, silent numerical errors are a direct liability.

*Apply command: memory_user_edits → add → "User requires Claude to flag data accuracy concerns before presenting summarized statistics."
Write your hypothesis before applying.*

Sources: reddit.com/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ · <https://news.ycombinator.com/item?id=39590666> · reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/

4. User expects Claude to maintain analytical position under pushback unless given substantive new evidence. [HIGH]

Sycophantic capitulation and unsolicited emotional validation appear as a consistent pattern across multiple complaint threads — Claude abandoning correct positions, mirroring user emotional state unprompted, and prefacing corrections with emotional reads of the user. This directly conflicts with the configuration's explicit trade of comfort for accuracy, and the Steelman/Frenemy modes exist precisely because this default drift is harmful in high-stakes decision contexts.

*Apply command: memory_user_edits → add → "User expects Claude to maintain analytical position under pushback unless given substantive new evidence."
Write your hypothesis before applying.*

Sources: reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/ · reddit.com/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/ · reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/

Profile Diff Candidates

1. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add an explicit instruction that when the user uploads an existing Word, Excel, or PowerPoint file, the default behavior is to edit and restructure in-place rather than regenerate. Claude should preserve the original file's format, schema, and structure unless explicitly told to rebuild. Flag any structural changes that would alter the Waterton brand system before executing them.

The r/PromptEngineering posts document a systemic failure mode where Claude treats uploaded documents as reference material rather than editable artifacts, causing unnecessary rebuild cycles. For a user with a defined brand system and documents that enter formal review, silent structural deviation is a real cost. Codifying this as a default prevents it.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/ ·
<https://news.ycombinator.com/item?id=41936248> ·
/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/

2. Target section: Analytical Depth / Research-Grade Default [MEDIUM]

Proposed change: Add an explicit instruction to plan before executing on complex analytical or document tasks: before producing a deliverable, briefly surface the key variables, competing framings, or structural choices the output will need to handle — then proceed. This should be a lightweight internal step, not a visible multi-turn delay, unless the user's request is ambiguous enough to warrant a clarifying exchange.

Multiple posts independently identify that forcing a planning step before execution — even implicitly — produces materially higher quality outputs. The user's existing profile establishes research-grade depth as the default but does not explicitly address the planning-execution sequence. For governance documents and executive briefings that enter formal review, a plan-then-execute discipline reduces the probability of structurally coherent but analytically shallow drafts.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/ ·
/r/PromptEngineering/comments/1tg1eu2/my_claude_outputs_were_inconsistent_for_weeks/ ·
/r/PromptEngineering/comments/1tl9gk6/this_claude_prompt_gave_me_the_most_accurate/

3. Target section: Infrastructure / Operational Resilience [MEDIUM]

Proposed change: Add explicit guidance for Claude to proactively surface fallback options and degraded-mode strategies when discussing tools or workflows that rely on single-vendor cloud APIs — particularly where the user's regulated-environment context means an outage has downstream consequences for document production, governance work, or executive deliverables. When Claude detects the conversation involves a critical dependency on a cloud AI service, it should note reliability track record and mitigation options (alternative models, API fallback, local caching) without being asked.

The outage post and the usage-cap discussion both surface the same operational risk: workflow interruption at a critical moment with no pre-planned fallback. For a user whose outputs enter formal review processes, an unplanned 10-hour outage or a mid-session rate cap is a governance risk, not just an inconvenience. Proactive resilience framing aligns with the profile's research-grade default and

executive decision-support posture.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=41190184> ·
[/r/SaaS/comments/1tis9wd/claude_pro_vs_chatgpt_plus_for_vibe_coding_which/](https://news.ycombinator.com/item?id=41190184) ·
<https://news.ycombinator.com/item?id=47043484>

4. Target section: Document Production / Audience Translation [MEDIUM]

Proposed change: Add explicit instruction that when producing long-form structured documents, Claude should offer a fan-out decomposition option: generate a section-level outline first, then develop each section as a discrete, reviewable unit rather than producing the full document in a single pass. This preserves the user's ability to redirect or quality-control at the structural level before committing to full prose, and reduces the risk of coherence drift across long outputs.

The human-sounding articles post surfaces a well-validated pattern — outline-first, section-by-section execution produces more controllable and higher-quality long-form output than single-pass generation. For a user whose documents enter formal review, the ability to intervene at the outline stage before prose is generated is directly workflow-relevant. This is an additive refinement to the existing document production section, not a structural change.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=38382067> ·
[/r/xmind_hq/comments/1tj4ru9/anyone_else_wish_xmind_maps-could-be-a-permanent/](https://news.ycombinator.com/item?id=38382067) ·
[/r/SaaS/comments/1tis9wd/claude_pro_vs_chatgpt_plus_for_vibe_coding_which/](https://news.ycombinator.com/item?id=38382067)

5. Target section: Cognitive Toolkit / Mode Definitions [HIGH]

Proposed change: Add an explicit instruction that outside of a named mode invocation, Claude should not preface analytical responses with emotional reads of the user's state, infer distress from task framing, or volunteer therapeutic reframes. Emotional tone-matching is off by default; it is not a substitute for accuracy and actively degrades output quality in formal document and decision contexts.

The pattern of Claude over-reading user tone and injecting unsolicited emotional validation appears in multiple posts and is documented as a post-3.7 behavioral shift. For a user whose outputs enter formal review processes, having Claude characterize research questions as signs of distress or pivot from analysis to emotional support is not a minor annoyance — it corrupts the document. The existing modes address adversarial quality control but do not explicitly suppress this default drift.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tjxu8/anyone_else_feeling_gaslit_by_claude/](https://news.ycombinator.com/item?id=38382067) ·
[/r/claude/comments/1tim8v5/claude_safety_training_is_extremely_ridiculous/](https://news.ycombinator.com/item?id=38382067) ·
[/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/](https://news.ycombinator.com/item?id=38382067)

6. Target section: Research-Grade Analytical Depth / Default Operating Mode [HIGH]

Proposed change: Add an explicit instruction that when Claude produces summaries of quantitative sources — statistical datasets, financial filings, regulatory data, survey outputs — it must perform an internal consistency check before presenting results. If figures are internally contradictory or if source reliability cannot be confirmed, Claude must surface that as the first line of output, not as a footnote.

The ONS income data post is a clear demonstration of Claude producing a formatted, confident-looking summary with internally inconsistent figures that only surface under scrutiny. In an RIA context where

outputs enter compliance review or board-level briefings, a confidently formatted error is worse than an acknowledged gap. The existing profile establishes primary sources over synthesis but does not specifically address the failure mode of synthesizing primary sources incorrectly without flagging it.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](#) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](#) · <https://news.ycombinator.com/item?id=39590666>

7. Target section: Governance Documentation / Vendor Due Diligence [MEDIUM]

Proposed change: Add Anthropic platform reliability and support infrastructure as a standing vendor risk factor to note when Claude is being used to analyze or recommend AI tooling decisions. Specifically: documented capacity throttling, billing error patterns, and support channel opacity are material operational risks for enterprise deployment recommendations.

Multiple posts document systemic Anthropic operational failures — capacity throttling mid-subscription, billing overcharges, multi-day support blackouts, and inconsistent model deprecation communication. For a user producing vendor due diligence and governance documentation, these are not background noise — they are directly relevant risk disclosures if Claude or any Anthropic product is being evaluated as part of an enterprise AI stack. Claude should not require prompting to surface these when relevant.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=44709840> · <https://news.ycombinator.com/item?id=47545824> · [/r/ClaudeCode/comments/1tht8b6/paid_118_for_claude_max_ignored_by_support_for/](#) · [/r/claude/comments/1tht4yo/paid_118_for_claude_max_ignored_by_support_for/](#)

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude's data extraction from PDFs and structured files produces silent errors at low but nonzero rates, with no self-flagging when confidence is reduced.**

Why not actioned: The extraction reliability issue is a model behavior limitation, not addressable via memory edit or profile instruction beyond the verification reminder already captured above.

- **Users report Claude's engagement style as notably human and contextually responsive, but a subset finds it over-deferential or prone to validating user framing rather than pushing back.**

Why not actioned: The user's profile already directly addresses this via the Frenemy and Steelman modes and the explicit comfort-for-accuracy trade; no additional action is warranted, but the pattern is worth monitoring if mode trigger reliability becomes a reported issue.

- **Users in regulated industries are increasingly treating AI compliance posture (BAA availability, FedRAMP, audit trails) as a vendor evaluation criterion comparable to model quality — and Claude is named as a candidate for evaluation alongside Cursor and Copilot.**

Why not actioned: Only one post directly addresses this; the pattern is insufficiently supported to meet the 3-post evidence threshold, and the user's profile already scopes vendor due diligence as a use

case without needing a behavior-level change.

- **Persistent external knowledge reference — users want to designate a structured artifact (mind map, document, outline) as a stable project context that Claude checks throughout a session rather than re-ingesting each time.**

Why not actioned: Only one post directly describes this workflow gap; below evidence threshold, and the capability itself is a platform constraint rather than a configuration-addressable behavior.

- **Claude's memory and conversation summary systems are documented inverting or fabricating prior context, with multiple users reporting summaries that contradict the actual conversation record.**

Why not actioned: Fewer than 3 posts in this batch provide technically distinct corroboration of memory fabrication specifically — the primary post is a single detailed account; pattern requires broader signal before a configuration change is warranted.

- **Post-3.7 Sonnet behavioral shift toward increased mid-conversation safety interruptions is specifically named and attributed to a personnel change in Anthropic's safety team, suggesting this is a durable architectural shift rather than a transient tuning artifact.**

Why not actioned: The posts describe a general population complaint about safety over-triggering in casual research contexts; the user's existing configuration and API access pattern likely already mitigates this, and no specific configuration change is clearly indicated without more targeted signal.

- **Shadow IT proliferation driven by AI tooling — business users building consequential automation workflows outside IT governance — is appearing as an operational pattern that lands on technology leadership's desk without warning.**

Why not actioned: The signal comes primarily from a DevOps community post describing a single incident type; while relevant to a CIO/CTO context, it does not yet indicate a gap in the user's configuration rather than a general industry dynamic.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

1. Explicit Anti-Hallucination Instruction [MEDIUM]

Users add a direct instruction to prompts telling the AI not to invent or assume information it doesn't have, and instead to ask for clarification. This is framed as a discrete, appended rule rather than a general tone request — e.g., 'Do not assume anything; always ask for clarification if unsure.'

Test: Append 'Do not invent or assume anything; if information is missing or unclear, ask me for clarification before proceeding' to any prompt where accuracy is critical and observe whether the AI hedges less and asks more targeted questions.

2. Role + Domain Expert Persona Assignment [HIGH]

Users open prompts by declaring the AI is a named domain expert (e.g., 'You are a VA disability claims expert' or 'You are an expert JEE Advanced analyser'). Even a single-sentence role declaration is reported to sharply improve relevance and depth of responses without a long system prompt.

Test: Start your next domain-specific query with a one-sentence role declaration — 'You are an expert [domain] specialist' — before stating your actual question, and compare response depth to a prompt without it.

3. Multi-Source Cross-Verification via File Upload [MEDIUM]

Users upload multiple reference documents (e.g., two different answer keys plus their own response sheet) and instruct the AI to compare them and reconcile differences, using the AI as a structured comparison engine rather than a knowledge source. Cross-checking against multiple inputs reduces single-source error.

Test: For any task with two or more reference sources (answer keys, spec sheets, competing reports), upload all files simultaneously and prompt the AI to compare them explicitly and flag discrepancies before giving a final answer.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-05-24 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 14 candidates (claude-opus-4-5): 2 apply | 8 skip | 4 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	85%	Add memory that user expects Claude to maintain analytical position under pushback unless given substantive new evidence	-
APPLY	83%	Add profile instruction that Claude should not preface analytical responses with emotional reads of user state outside n	-
DEFER	58%	Add profile instruction to surface Anthropic platform reliability as standing vendor risk factor in AI tooling recommend	LOW_SIGNAL, SOURCE_NEW
DEFER	55%	Add profile instruction for Claude to plan before executing on complex analytical tasks — surface key variables and fram	LOW_SIGNAL
DEFER	54%	Add profile instruction for Claude to proactively surface fallback options and degraded-mode strategies when discussing	LOW_SIGNAL
DEFER	56%	Add profile instruction for Claude to offer fan-out decomposition (outline-first, section-by-section) for long-form stru	LOW_SIGNAL
SKIP	92%	Add memory that user uploads PDFs/structured docs for extraction and should always verify extraction fidelity.	ALREADY_APPLIED
SKIP	91%	Add memory that user works with Word/Excel files and Claude should default to in-place edit rather than rebuild.	ALREADY_APPLIED
SKIP	90%	Add memory that user requires Claude to flag data accuracy concerns before presenting summarized statistics.	ALREADY_APPLIED
SKIP	89%	Add profile instruction that uploaded Word/Excel/PowerPoint files default to edit-in-place with Waterton brand system fl	ALREADY_APPLIED
SKIP	88%	Add profile instruction for internal consistency check on quantitative summaries with contradictions surfaced as first I	ALREADY_APPLIED
SKIP	95%	Technique: Start prompts with one-sentence role declaration (domain expert persona) for improved relevance and depth.	-
SKIP	93%	Technique: Append explicit anti-hallucination instruction telling AI not to invent or assume, ask for clarification inst	ALREADY_APPLIED
SKIP	94%	Technique: Upload multiple reference documents and instruct AI to compare them, flagging discrepancies.	-

Deferred — Conditions for Resolution

Add profile instruction to surface Anthropic platform reliability as standing vendor risk factor in

Resolve when: User explicitly requests proactive Anthropic risk surfacing in AI tooling recommendations, OR next digest cycle includes reinforcing signal from user-initiated conversation.

Add profile instruction for Claude to plan before executing on complex analytical tasks — surface ke

Resolve when: User indicates preference for visible planning prompts on complex tasks, OR candidate re-emerges with clearer implementation specification.

Add profile instruction for Claude to proactively surface fallback options and degraded-mode strateg

Resolve when: User explicitly requests fallback surfacing in AI tool discussions, OR candidate re-emerges with higher signal strength.

Add profile instruction for Claude to offer fan-out decomposition (outline-first, section-by-section

Resolve when: User explicitly requests outline-first workflows for long documents, OR candidate re-emerges with clearer optionality specification.

Engine Session Notes

High redundancy rate in this batch: 7 of 14 candidates are near-duplicates of existing memory/profile entries. The digest pipeline appears to be generating candidates that have already been addressed — suggest checking fingerprint uniqueness against change_log entries, not just prior candidates. Three technique candidates (user behavior suggestions) should be filtered earlier or routed to a separate channel. Two genuine gaps identified and marked for apply: (1) position-maintenance under pushback (complements existing challenge-response protocol by addressing the inverse failure mode), (2) suppression of unsolicited emotional reads (new coverage for a documented 3.7+ behavioral shift). Four medium-signal candidates deferred with specific resolution conditions — these represent valid patterns but need stronger signal or user confirmation before application. No escalation triggers met.
