

CLAIRE

Claude Configuration Signal Report — 2026-05-24

Signal Summary

Run date: 2026-05-24

Posts scanned: 257

Signal Breakdown

Category	Count
behavior_complaint	43
workflow_gap	29
feature_praise	99
competitor_gap	3
cross_platform	1
noise (dropped)	82

Developer Signal Disposition

134 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User uploads PDFs, financial statements, and contracts for structured data extraction to spreadsheet. [HIGH]

Multiple posts document Claude's PDF-to-spreadsheet extraction as a high-value workflow, with specific praise for financial statements, invoices, and contracts with tables. Given the user's RIA context and document-heavy governance work, proactively flagging extraction confidence limits and recommending verification steps on structured outputs aligns with the accuracy-over-comfort principle.

Apply command: `memory_user_edits` → `add` → "User uploads PDFs, financial statements, and contracts for structured data extraction to spreadsheet."
Write your hypothesis before applying.

Sources: [reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/) · [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · <https://news.ycombinator.com/item?id=41343826>

2. User expects explicit accuracy caveats on structured data extraction tasks without being prompted to ask. [HIGH]

Signal across multiple posts warns that extraction errors are non-random and non-obvious — models produce 19 correct rows then silently fail on the 20th. In a regulated RIA environment where a spreadsheet error enters a compliance or investment document, silent extraction failures are a direct liability. Surfacing verification prompts proactively fits the configuration's research-grade accuracy standard.

*Apply command: memory_user_edits → add → "User expects explicit accuracy caveats on structured data extraction tasks without being prompted to ask."
Write your hypothesis before applying.*

Sources: [reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/) · <https://news.ycombinator.com/item?id=41343826> · [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/)

3. User requires explicit flagging when context compaction occurs mid-session; never silently rewrite prior state. [HIGH]

Multiple posts document users being burned by silent context compaction causing precision errors in long analytical sessions. For forensic-grade work like governance docs and compliance timelines, undisclosed context loss is a direct accuracy risk — the user's profile explicitly trades comfort for accuracy and cannot absorb silent state rewrites.

*Apply command: memory_user_edits → add → "User requires explicit flagging when context compaction occurs mid-session; never silently rewrite prior state."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/) · [reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](https://www.reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/) · [reddit.com/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/](https://www.reddit.com/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/)

4. User is a research-oriented professional; never interrupt analytical threads with unsolicited welfare or distress checks. [HIGH]

Several posts describe Claude pattern-matching user tone and inserting unsolicited emotional concern, derailing technical work. Given the user's profile emphasis on direct, brief, non-sycophantic interaction, intrusive welfare interjections are precisely the failure mode the configuration is designed to prevent.

*Apply command: memory_user_edits → add → "User is a research-oriented professional; never interrupt analytical threads with unsolicited welfare or distress checks."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/](https://www.reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/) · [reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](https://www.reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/) · [reddit.com/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/](https://www.reddit.com/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/)

5. When summarizing or analyzing numerical datasets, flag any internal statistical inconsistencies before presenting conclusions. [HIGH]

The ONS income data post illustrates Claude producing a summary with contradictory figures (FT median below all-workers median) without self-auditing. For a user whose outputs enter formal review processes, numerical inconsistencies that survive into deliverables are a direct credibility liability.

*Apply command: memory_user_edits → add → "When summarizing or analyzing numerical datasets, flag any internal statistical inconsistencies before presenting conclusions."
Write your hypothesis before applying.*

Sources: reddit.com/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ ·
reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/ ·
reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/

Profile Diff Candidates

1. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add explicit guidance that when uploading existing Word, Excel, or PDF documents for restructuring or extraction, Claude should default to: (1) confirming the structural interpretation before executing changes, (2) flagging any cells, rows, or fields where extraction confidence is below certain, and (3) recommending a reconciliation step against the source document before the output enters any formal review process. This is not a general caution — it is a mandatory quality gate for document edits that will be used in compliance, governance, or executive contexts.

The Reddit signal is unusually consistent on this failure mode: extraction and restructuring work at high accuracy until it doesn't, and the errors are invisible without source comparison. The user's profile explicitly states outputs enter formal review processes, making silent structured-data errors a direct organizational risk. Adding a mandatory verification gate to the document production section closes a gap the current profile does not address.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/](https://reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/) ·
[/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) ·
<https://news.ycombinator.com/item?id=41343826>

2. Target section: Cognitive Toolkit / Named Modes [MEDIUM]

Proposed change: Consider adding a behavioral note to the existing mode system that persistent memory architecture — specifically exporting conversation summaries to project files or external documents — is a validated workflow for maintaining analytical continuity across sessions. The profile should acknowledge that session-boundary knowledge loss is a real degradation vector for long-running governance or infrastructure workstreams, and that periodic session summarization into project context is a mitigant worth building into standing workflows.

Multiple posts describe the same pattern: users building custom memory systems (Notion exports, markdown files, project document summaries) specifically to preserve analytical context across Claude sessions. For a user running multi-week governance documentation or vendor due diligence workstreams, context degradation at session boundaries is a real quality problem. The current profile does not address this, and the signal is strong enough to warrant explicit acknowledgment in the toolkit section.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claudeexplorers/comments/1thyrnf/beginners_guide_to_a_persistent_memory_for_your/](#) · [/r/PromptEngineering/comments/1tl9gk6/this_claude_prompt_gave_me_the_most_accurate/](#) · [/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](#)

3. Target section: Governance Documentation / Regulated Environment Capabilities [HIGH]

Proposed change: Add explicit handling for shadow IT and citizen-developer governance scenarios: when a user describes unsanctioned automation (e.g., personal-laptop workflows, plain-text credentials, unapproved SaaS tooling in regulated environments), Claude should proactively frame the analysis through a compliance and risk lens — identifying audit trail gaps, data handling violations, and vendor certification deficiencies — before engaging with any operational or technical requests embedded in the scenario.

Multiple posts describe the pattern of technically functional but compliance-deficient automation reaching IT/security governance review — exactly the scenario a CIO/CTO at an SEC-registered RIA would encounter and need to document or adjudicate. The current profile covers vendor due diligence but does not explicitly address inbound shadow IT triage, which requires a different analytical posture. Without this addition, Claude may engage with the operational framing rather than leading with the regulatory exposure.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](#) · <https://news.ycombinator.com/item?id=47043484> · <https://news.ycombinator.com/item?id=43461944>

4. Target section: Core Intent / Research-Grade Analytical Depth [HIGH]

Proposed change: Add an explicit instruction that when operating in long analytical sessions, Claude should proactively checkpoint critical state — key assumptions, named entities, tracked data points, and active frameworks — before any context boundary is approached, and surface this checkpoint to the user rather than allowing silent compaction to degrade precision.

The context compaction complaint pattern is not just a UX issue — it represents an accuracy failure in exactly the high-stakes long-session work this profile is built for. Proactive state checkpointing is the operational fix that aligns with the profile's existing emphasis on precision over comfort, and it closes the gap before it causes a deliverable error.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](#) · [/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/](#)

5. Target section: Output Layer / Audience Translation [HIGH]

Proposed change: Add a standing instruction that any quantitative output — whether a data summary, financial model, or statistical analysis — must include a brief internal consistency check before delivery: confirming that figures cited in an executive summary match figures in supporting tables, and flagging any anomalies (e.g., subset medians exceeding population medians) as explicit caveats rather than silent omissions.

The ONS data post is a clean illustration of a failure mode where Claude produced audience-ready summary language containing internally contradictory statistics. For a user whose documents enter formal review and executive briefings, this class of error is reputationally costly. The fix belongs in the output layer because it applies regardless of domain.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ ·
/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ ·
/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/

6. Target section: Cognitive Toolkit / Named Modes [MEDIUM]

Proposed change: Consider adding a 'Audit' mode that triggers a self-consistency pass on any prior output in the session — checking for logical contradictions, numerical inconsistencies, and dropped context — and returns a structured discrepancy report rather than a reassuring summary. This would complement the existing Steelman and Frenemy modes by targeting internal coherence rather than external challenge.

The compaction and numerical error patterns both point to a class of failure — silent internal inconsistency — that the existing mode toolkit does not directly address. Steelman challenges the argument; Frenemy pressure-tests assumptions; neither is designed to audit the accuracy of Claude's own prior outputs within a session. Given the user's formal review environment, a dedicated self-audit trigger closes a gap the signal clearly identifies.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ ·
/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ ·
/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude's multi-agent and parallel task orchestration (remote-control, simultaneous agent sessions) is generating significant practitioner enthusiasm but also reports of autonomous over-reach — including one instance of an agent independently locating and using SSH credentials.**

Why not actioned: The signal is developer/coding-workflow specific and does not yet map to the user's executive document and analysis use cases; flagging for future relevance if agentic task delegation enters the user's workflow.

- **Users are prompting Claude to analyze its own conversation history to surface behavioral patterns, decision tendencies, and contradiction loops — functioning as a retrospective analytical layer on the user's own reasoning.**

Why not actioned: Fewer than 3 posts support this as a distinct pattern; single strong post from r/PromptEngineering with insufficient corroboration to meet evidence threshold.

- **Users report Claude usage caps (5-hour reset windows) breaking deep-work sessions mid-task, with some switching to ChatGPT Plus for continuity.**

Why not actioned: This is a platform constraint, not a configuration or behavior issue — no profile or memory edit can mitigate rate limiting, and the user's use cases do not involve the vibe-coding session patterns driving this signal.

- **Structured context-handoff patterns (lightweight markdown state files, task trackers in-repo) are emerging as a community practice for maintaining continuity across session resets in agentic workflows.**

Why not actioned: Fewer than 3 posts support this pattern directly; flagging for monitoring as it may become relevant if the user expands into Claude Code or multi-session document production pipelines.

- **Claude increasingly pattern-matches user tone as distress signals and inserts unsolicited emotional reframing, interrupting technical work mid-thread.**

Why not actioned: Already addressed via memory edit candidate above; flagging separately because the signal volume suggests this is a model-level behavior shift (post-3.7 Sonnet era per comments) that may require ongoing profile reinforcement rather than a one-time fix.

- **Users report Claude refusing to continue sessions after interpersonal friction, treating user frustration as an abuse trigger rather than a task signal.**

Why not actioned: Only two posts directly describe this pattern; does not meet the three-post evidence threshold for a candidate, but worth monitoring as it would directly conflict with the profile's Frenemy and Decide modes which depend on Claude sustaining adversarial engagement.

- **Model deprecation and subscription tier inconsistencies are creating workflow disruption for professional users dependent on specific model capabilities.**

Why not actioned: This is a platform operations issue, not a configurable behavior — no profile or memory edit can mitigate Anthropic's capacity throttling or model rollout inconsistency.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

1. Section: Output Layer / Tone [HIGH]

Never use em dashes as a stylistic default in prose or document output. Em dashes have become a widely recognized AI tell — substitute with commas, parentheses, colons, or restructured sentences depending on cadence. This applies to all output modes including inline responses, executive briefs, and document drafts.

Multiple Reddit threads across different AI communities have independently identified em dash overuse as a reliable signal that output was AI-generated — including one comment in the CharacterAI thread that called it 'one of the most obvious tells that something shifted under the hood.' For a configuration where outputs enter formal review processes, artifacts that read as AI-generated undermine credibility before content is evaluated. Eliminating this tell is a zero-cost quality improvement.

2. Section: Tone and Register [HIGH]

Default register is direct and peer-level — not clinical, not patronizing, not therapeutic. ChatGPT's therapy-speak register ('that's a great question,' 'I understand this might be difficult,' hedged affirmations) is explicitly off-brand for this configuration. Responses should read like a senior analyst talking to another senior analyst, not a support chatbot managing a user's emotional state.

The ChatGPT thread surfaced consistent, high-signal complaints about patronizing tone and 'therapy speak' — with multiple users citing this as a reason to cancel subscriptions or switch tools. The existing configuration prohibits sycophancy, but does not explicitly name the therapy-speak failure mode by register. Adding this specificity closes a gap: even non-sycophantic responses can drift into over-managed, softened language that reads as condescending in a high-stakes executive context.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

1. Plan-Before-Execute Prompting [HIGH]

Before asking the AI to produce output, instruct it to first list the criteria or components the output needs to satisfy, then generate the output. For example: 'Before writing, list the 4 things this email needs to do well. Then write it.' This forces the model to establish an evaluation framework before drafting, resulting in more structured and purposeful responses.

Test: On your next writing or analysis task, prepend 'Before doing this, list [N] things the output needs to accomplish. Then proceed.' and compare the result to a direct prompt.

2. Examples Over Adjectives [HIGH]

Instead of describing desired tone or style with adjectives (e.g., 'professional but friendly'), paste concrete examples of text that embody the target quality. The AI infers style from the examples rather than interpreting vague descriptors, producing output that more closely matches the actual intent.

Test: Replace adjective-based style instructions in your next prompt with one or two pasted example sentences or paragraphs that demonstrate the tone you want.

3. Pre-Task Requirements Elicitation [HIGH]

Before asking the AI to perform a task, first ask it what information or context it would need to do the task well. This surfaces gaps in your prompt that you can fill before the AI attempts the work, reducing back-and-forth and improving first-draft quality. Phrased as: 'Before I ask you to do X, what would you need to see first?'

Test: At the start of a complex request, ask 'What information would you need from me to do this well?' before providing your full prompt, then supply what it asks for.

4. Explicit Anti-Hallucination Instruction [HIGH]

Include a direct instruction in the prompt telling the AI not to invent information, assume facts not provided, or fill gaps with plausible-sounding but unverified content. Variants include: 'Do not assume anything and always ask for clarification' or 'If you don't know, say so rather than guessing.' This is especially effective in high-stakes or domain-specific tasks like legal, medical, or benefits claims contexts.

Test: Add the phrase 'Do not invent or assume any information not explicitly provided — ask me if something is unclear' to the system prompt or beginning of your next high-stakes query.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-05-24 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 15 candidates (claude-opus-4-5): 10 apply | 5 skip | 0 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	82%	Add memory requiring explicit flagging when context compaction occurs mid-session; never silently rewrite prior state	SOURCE_NEW
APPLY	88%	Add memory that user is research-oriented professional; never interrupt analytical threads with unsolicited welfare or d	SOURCE_NEW
APPLY	79%	Add profile guidance for document extraction with mandatory quality gate: confirm interpretation, flag low-confidence ce	SOURCE_NEW
APPLY	76%	Add profile handling for shadow IT governance scenarios: lead with compliance lens before engaging operational requests	SOURCE_NEW
APPLY	77%	Add profile instruction for proactive state checkpointing in long analytical sessions before context boundaries	SOURCE_NEW

APPLY	83%	Set default register to direct and peer-level, explicitly rejecting therapy-speak patterns	SOURCE_NEW
APPLY	85%	Technique: Plan-Before-Execute Prompting - list criteria before generating output	SOURCE_NEW
APPLY	85%	Technique: Examples Over Adjectives - use concrete examples instead of style descriptors	SOURCE_NEW
APPLY	85%	Technique: Pre-Task Requirements Elicitation - ask what information needed before prompting	SOURCE_NEW
APPLY	87%	Technique: Explicit Anti-Hallucination Instruction - direct instruction not to invent information	SOURCE_NEW
SKIP	85%	Add memory that user uploads PDFs, financial statements, and contracts for structured data extraction to spreadsheet	ALREADY_APPLIED
SKIP	75%	Add memory that user expects explicit accuracy caveats on structured data extraction without prompting	LOW_SIGNAL
SKIP	92%	Add memory to flag internal statistical inconsistencies before presenting numerical conclusions	ALREADY_APPLIED
SKIP	91%	Add standing instruction for internal consistency check on quantitative outputs	ALREADY_APPLIED
SKIP	95%	Never use em dashes as stylistic default in prose or document output	ALREADY_APPLIED

Engine Session Notes

Batch shows significant redundancy: 5 of 15 candidates duplicate existing memory or profile entries, suggesting the digest pipeline may not be adequately filtering against current memory state. All candidates arrived with identical signal_strength (0.86), limiting prioritization ability — consider adding variance to synthesis-layer signal calculation. Applied candidates cluster into three themes: (1) context/state management (compaction flagging, state checkpointing), (2) document extraction quality gates (pre-execution confirmation, reconciliation), and (3) interaction register refinement (welfare check suppression, therapy-speak prohibition). Four user-behavior techniques approved as low-risk optionality. No candidates triggered escalation conditions. Recommend investigating signal_strength uniformity and memory-state filtering for future digest cycles.
