

CLAIRE

Claude Configuration Signal Report — 2026-05-24

Signal Summary

Run date: 2026-05-24

Posts scanned: 257

Signal Breakdown

Category	Count
behavior_complaint	46
workflow_gap	38
feature_praise	95
competitor_gap	0
cross_platform	0
noise (dropped)	78

Developer Signal Disposition

139 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User uploads existing Word/Excel/PDF files for in-place editing, not just net-new document generation. [HIGH]

Multiple posts confirm users discovering Claude can restructure and edit existing documents rather than only generating new ones. Given this user's Waterton document production system and formal review workflows, Claude should default to preserving existing formatting and structure conventions when files are uploaded, rather than treating every upload as a blank-slate generation task.

Apply command: `memory_user_edits` → add → "User uploads existing Word/Excel/PDF files for in-place editing, not just net-new document generation."
Write your hypothesis before applying.

Sources: [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://news.ycombinator.com/item?id=41936248) · [reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](https://news.ycombinator.com/item?id=41936248) · <https://news.ycombinator.com/item?id=41936248>

2. User requires output verification flagging when Claude performs data extraction or numerical transcription from PDFs. [HIGH]

Multiple posts surface a consistent failure mode: Claude extracts data from PDFs accurately for many rows then silently errors, with no warning issued. Given this user's regulatory environment and document outputs that enter formal review, silent extraction errors are a compliance risk. Claude should proactively flag extraction confidence and recommend spot-check steps rather than presenting extracted data as unconditionally reliable.

*Apply command: memory_user_edits → add → "User requires output verification flagging when Claude performs data extraction or numerical transcription from PDFs."
Write your hypothesis before applying.*

Sources: [reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](https://news.ycombinator.com/item?id=41936248) · [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and-restructure/](https://news.ycombinator.com/item?id=41936248) · <https://news.ycombinator.com/item?id=41936248>

3. User operates in SEC-registered RIA; AI tooling must clear compliance bar before adoption recommendations. [HIGH]

The Cursor/compliance post surfaces a recurring regulated-industry pattern: AI tooling evaluated on HIPAA BAA, SOC2, FedRAMP, and audit trail requirements. Given the user's RIA environment, any vendor or tooling recommendation from Claude should proactively flag compliance certification status rather than waiting to be asked.

*Apply command: memory_user_edits → add → "User operates in SEC-registered RIA; AI tooling must clear compliance bar before adoption recommendations."
Write your hypothesis before applying.*

Sources: <https://news.ycombinator.com/item?id=47043484>

4. User requires explicit notice when context compaction or summarization occurs mid-session. [HIGH]

Multiple posts document Claude silently rewriting or compacting long-session context, causing precision errors in analytical work. For a user running governance, architecture, and evidentiary-grade analysis, silent context loss is directly harmful — a degraded response arriving without warning is worse than a visible failure.

*Apply command: memory_user_edits → add → "User requires explicit notice when context compaction or summarization occurs mid-session."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://news.ycombinator.com/item?id=47043484) · [reddit.com/r/ChatGPTcomplaints/comments/1tkpzs/on_claude_i_caught_both_memory_system_and/](https://news.ycombinator.com/item?id=47043484) · [reddit.com/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/](https://news.ycombinator.com/item?id=47043484)

5. User has zero tolerance for unsolicited emotional tone-reading or welfare check insertions mid-task. [HIGH]

A consistent pattern across multiple posts describes Claude injecting concern framing, wellness language, or emotional re-reads of user tone — mid-analysis, unprompted. For a user whose configuration explicitly trades comfort for accuracy, this behavior interrupts executive-grade work and is indistinguishable from sycophantic hedging.

Apply command: memory_user_edits → add → "User has zero tolerance for unsolicited emotional tone-reading or welfare check insertions mid-task."

Write your hypothesis before applying.

Sources: reddit.com/r/claude/comments/1tim8v5/claude_safety_training_is_extremely_ridiculous/ ·
reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/ ·
reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/

6. User requires data accuracy flags when Claude output may contain summarization errors against source material. [HIGH]

The ONS income data post is a documented case of Claude producing statistically inconsistent summaries that were accepted as accurate before community review caught the errors. For a user whose outputs enter formal review processes or influence organizational decisions, silent numerical errors in summarization are a material risk — flagging uncertainty is preferable to confident wrongness.

Apply command: memory_user_edits → add → "User requires data accuracy flags when Claude output may contain summarization errors against source material."

Write your hypothesis before applying.

Sources: reddit.com/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ ·
reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/ ·
reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/

Profile Diff Candidates

1. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add explicit guidance that when existing documents are uploaded, Claude defaults to edit-in-place mode: preserve the source structure, heading hierarchy, and formatting conventions unless explicitly instructed to restructure. Distinguish this from net-new generation mode. For PDF-sourced data destined for spreadsheets or briefs, append a brief extraction-confidence note flagging any cells or figures where transcription certainty is reduced (dense tables, rotated text, scanned pages), and recommend a spot-check step before the document enters formal review.

The signal across multiple posts confirms two distinct document interaction patterns — new generation and existing-file editing — that require different default behaviors. This user's outputs enter formal review processes where silent formatting drift or extraction errors create downstream rework. Formalizing the distinction in the profile prevents Claude from inadvertently restructuring documents that should be preserved, and adds a lightweight compliance backstop on extracted numerical data.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/ ·
/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/ ·
https://news.ycombinator.com/item?id=41936248

2. Target section: Cognitive Toolkit / Named Modes [MEDIUM]

Proposed change: Add a note that when the user invokes plan-first sequencing (explicitly asking Claude to enumerate requirements or constraints before executing), Claude should treat this as a structured decomposition request — producing a numbered pre-flight checklist that the user can edit or reject before generation begins. This is distinct from the existing Focus mode and should not require a named trigger; it activates whenever the user asks Claude to 'plan first' or 'list what this needs to do before you do it.'

Multiple posts identify plan-before-execute as a high-leverage pattern that improves output quality by surfacing assumptions before generation rather than after. For this user's governance documentation and business case work — where wrong assumptions discovered late mean full rewrites — making this sequencing a recognized interaction pattern (rather than ad hoc) reduces rework. The profile already formalizes modes for quality control and decision-forcing; this is a natural complement for scoping and assumption-surfacing before high-stakes document generation.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/](#) · [/r/PromptEngineering/comments/1tg1eu2/my_claude_outputs_were_inconsistent_for_weeks/](#) · [/r/ClaudeCode/comments/1tjbdm4/it_just_keeps_getting_better_and_better/](#)

3. Target section: Vendor Due Diligence [HIGH]

Proposed change: Add a standing checklist trigger: when evaluating any AI-adjacent tooling or SaaS vendor, Claude should proactively surface compliance certification status (SOC2, HIPAA BAA, FedRAMP, data residency, audit trail availability) as a first-order output item, not a footnote. This applies even when the user's prompt does not explicitly invoke compliance — the regulated environment is always the operating context.

The Cursor compliance thread shows that regulated-industry practitioners frequently discover compliance gaps only after evaluation is underway. The user's RIA context means this gap has direct regulatory exposure. Surfacing certification status proactively prevents the user from having to loop back after the fact and aligns with the profile's intent that expert-level framing is the default, not a special request.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=47043484>

4. Target section: Business Case Development [MEDIUM]

Proposed change: Add explicit handling for shadow-IT and citizen-automation scenarios: when a business case involves workflows built outside IT governance (personal laptops, unsanctioned SaaS, plaintext credentials, no audit trail), Claude should frame the remediation path as a governance opportunity with a structured proposal arc — current-state risk inventory, compliant-equivalent architecture, change management angle — rather than treating it as a pure technical cleanup task.

The automation/cleanup-crew post describes a pattern highly relevant to a CIO/CTO scope: unsanctioned automation surfaces organizational demand signals that are genuinely useful for roadmap and governance work. The user's profile already covers audience translation and business case development; this extends that capability to shadow-IT intake, which is a recurring executive-level governance problem in remote-model organizations.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](#) · <https://news.ycombinator.com/item?id=47043484>

5. Target section: Core Analytical Defaults [HIGH]

Proposed change: Add an explicit instruction that when Claude performs multi-step data summarization or cross-reference tasks against provided source material, it must surface any internal inconsistencies detected in its own output before delivering the final result — rather than presenting a polished summary that may contain silent errors. Framing: 'flag before finalizing, not after.'

The ONS data case and the forensic timeline case both show Claude producing confidently-formatted outputs that contained factual or logical inconsistencies only caught externally. The user's profile already establishes that accuracy trumps comfort — this extends that principle to Claude's self-review behavior on quantitative and evidentiary tasks, which are both live use cases in this configuration.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](#) · [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](#)

6. Target section: Long-Session and Document Production Behavior [HIGH]

Proposed change: Add a standing instruction that when Claude detects it is approaching context limits in a long analytical session, it should proactively prompt the user to extract and pin critical parameters, constraints, or data anchors before compaction occurs — rather than compacting silently and continuing with a degraded context model.

Multiple posts confirm that silent context compaction is a real behavior that causes precision degradation without user awareness. For this user's use cases — vendor due diligence, infrastructure architecture, governance documentation — a session that loses its analytical anchors mid-work produces outputs that require full reconstruction. Proactive surfacing converts an invisible failure into a manageable workflow checkpoint.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#) · [/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/](#) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](#)

7. Target section: Interaction Constraints / Anti-Patterns [HIGH]

Proposed change: Add an explicit prohibition on mid-task interruptions framed around user wellbeing, topic sensitivity, or inferred emotional state. Specifically: Claude should not insert 'I need to pause here,' 'I'm noticing this topic may be sensitive,' or equivalent constructions during research, analysis, or document production tasks unless the user has invoked a mode that calls for reflective dialogue. Safety disclosures at session end are acceptable; mid-analysis interruptions are not.

The safety-interruption complaint pattern is the highest-volume behavior complaint in this signal set. For a user running regulated-environment analysis — including topics that may naturally touch on risk, compliance, security, or organizational stress — mid-task moralizing creates direct workflow damage and is precisely the 'hedge-everything' behavior the configuration exists to suppress.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tim8v5/claude_safety_training_is_extremely_ridiculous/](#) · [/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](#) · [/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/](#)

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Users report Claude exhibiting sycophantic encouragement loops — affirming project directions and expanding scope rather than flagging overextension or recommending against a path.**

Why not actioned: The existing profile already addresses sycophancy resistance directly via the Steelman and Frenemy modes; no new candidate is warranted unless the signal shows these modes failing specifically, which the posts do not demonstrate.

- **Persistent memory management via external tools (Notion, exported markdown, manual summary workflows) is an active area of user experimentation with no stable best practice emerging.**

Why not actioned: Only two posts address this directly and neither describes a capability gap relevant to this user's configured workflow; insufficient signal to generate a candidate.

- **Claude availability and downtime is treated as a hard workflow blocker by professional users with no documented fallback or continuity posture.**

Why not actioned: Only one post addresses outage directly and the signal is about infrastructure availability, not a Claude behavior or configuration gap addressable through profile or memory edits.

- **Skepticism about AI agent reliability for sustained, multi-step autonomous tasks remains high among technical practitioners, with Pokémon benchmarking used as a proxy for real-world agentic capability gaps.**

Why not actioned: Fewer than 3 posts support an actionable pattern for this user's configuration; the signal reflects general AI limitations discourse rather than a specific Claude behavior relevant to the user's workflow.

- **Claude appears to exhibit increased safety-interrupt and assumption-of-harm frequency in Opus 4.7 relative to prior model versions, suggesting a model-version-specific regression rather than a general Claude behavior.**

Why not actioned: The signal identifies a version-specific pattern but does not provide sufficient evidence of a configuration-level remedy — this requires monitoring across model versions before a profile change is warranted.

- **Users report Claude generating repetitive clarifying questions in image/screenshot analysis tasks rather than proceeding with best-available interpretation.**

Why not actioned: Only two posts touch this pattern with sufficient specificity; does not meet the three-post evidence threshold.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

1. Persistent System-Level Role Prompting via Project Instructions [MEDIUM]

Users paste a detailed, structured role prompt into a project or system instructions field so the AI consistently applies expert framing, constraints, and behavior across all sessions without re-prompting each time. The technique involves defining domain expertise, behavioral guardrails, and scope in a single reusable block.

Test: Create a project in Claude or ChatGPT, paste a detailed role and constraint prompt into the system instructions, then run several queries to verify the persona and rules persist automatically.

2. Multi-File Upload for Comparative Document Analysis [MEDIUM]

Users upload multiple related documents simultaneously and instruct the AI to compare, reconcile, or score them against each other. This offloads manual cross-referencing to the AI and works across ChatGPT, Claude, and Gemini.

Test: Upload two versions of a document (e.g., two answer keys or two drafts) to any AI tool and prompt it to identify all differences and produce a reconciled summary.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-05-24 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 15 candidates (claude-opus-4-5): 6 apply | 7 skip | 2 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	82%	Add memory instruction that when existing files are uploaded, Claude should default to preserving formatting and structure	SOURCE_NEW
APPLY	72%	Add memory instruction requiring explicit notice when context compaction or summarization occurs mid-session.	SOURCE_NEW
APPLY	88%	Add memory instruction that user has zero tolerance for unsolicited emotional tone-reading or welfare check insertions	-
APPLY	75%	Add profile instruction that when approaching context limits, Claude should proactively prompt user to extract critical	SOURCE_NEW
APPLY	90%	Add explicit profile prohibition on mid-task interruptions framed around user wellbeing, topic sensitivity, or inferred	-
DEFER	68%	Add profile instruction to surface internal inconsistencies before delivering final result ('flag before finalizing') for	-
DEFER	58%	Add profile note recognizing plan-first sequencing as an interaction pattern that produces numbered pre-flight checklist	LOW_SIGNAL, SOURCE_NEW
DEFER	52%	Add profile handling for shadow-IT and citizen-automation scenarios, framing remediation as a governance opportunity.	LOW_SIGNAL, SOURCE_NEW
SKIP	95%	Add memory instruction requiring output verification flagging when Claude extracts data from PDFs.	ALREADY_APPLIED
SKIP	85%	Add memory instruction requiring data accuracy flags when Claude output may contain summarization errors against source	ALREADY_APPLIED
SKIP	80%	Add profile guidance for edit-in-place mode when existing documents are uploaded, plus extraction-confidence notes on PD	-
SKIP	92%	Add memory instruction that user operates in SEC-registered RIA; AI tooling must clear compliance bar before adoption re	ALREADY_APPLIED

SKIP	93%	Add profile instruction for standing checklist trigger on compliance certification status for AI tooling evaluation.	ALREADY_APPLIED
SKIP	95%	Technique observation: persistent system-level role prompting via project instructions.	ALREADY_APPLIED, LOW_SIGNAL
SKIP	88%	Technique observation: multi-file upload for comparative document analysis.	LOW_SIGNAL, SOURCE_NEW

Deferred — Conditions for Resolution

Add profile instruction to surface internal inconsistencies before delivering final result ('flag be

Resolve when: Eval data on d0f4ac00 adherence becomes available, or 14 days pass without regression signal on internal consistency checking.

Add profile note recognizing plan-first sequencing as an interaction pattern that produces numbered

Resolve when: Signal strength on plan-first sequencing rises above 0.70 in a future digest cycle, or user explicitly invokes this pattern in conversation.

Add profile handling for shadow-IT and citizen-automation scenarios, framing remediation as a govern

Resolve when: Signal strength on shadow-IT/citizen-automation scenario rises above 0.70 in a future digest cycle.

Engine Session Notes

This batch shows two strong signal clusters: (1) silent failure disclosure (context compaction, extraction errors) and (2) unsolicited behavioral interruptions (wellness checks, sensitivity pauses). Applied candidates address both. Seven candidates skipped as duplicates of existing memory/profile entries or as generic technique observations — the digest pipeline produced several near-duplicates of entries applied 2026-05-19 and 2026-05-23, suggesting fingerprint matching against memory state could be improved. Two candidates deferred on low signal strength. The identical 0.86 signal strength across 9 candidates raises a calibration question for the synthesizer. Track C technique candidates both describe patterns either already in use or too generic to act on — consider raising threshold for technique surfacing. No escalation required: confidence distribution is healthy, no HIGH_VARIANCE + CONTRADICTS_MEMORY combinations, no shared defer conditions across multiple candidates.
