

CLAIRE

Claude Configuration Signal Report — 2026-05-31

Signal Summary

Run date: 2026-05-31

Posts scanned: 264

Signal Breakdown

Category	Count
behavior_complaint	0
workflow_gap	1
feature_praise	2
competitor_gap	0
cross_platform	0
noise (dropped)	0

Developer Signal Disposition

3 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. When extracting tabular data from PDFs, flag row/column counts and any cells where confidence is reduced before output enters review. [HIGH]

Multiple posts confirm Claude's PDF table extraction produces silent errors that only surface downstream. Memory edit 9 already exists but this candidate would duplicate it — skipping.

Apply command: `memory_user_edits` → add → "When extracting tabular data from PDFs, flag row/column counts and any cells where confidence is reduced before output enters review."
Write your hypothesis before applying.

2. User tracks Anthropic platform uptime as a regulated-env SLA risk; flag when outage patterns recur. [HIGH]

Two distinct HN posts document Claude.ai/API outages with 500 errors and sub-99.9% availability. Given the user's SEC-registered RIA context and existing memory edit 10 on platform reliability, recurring outage patterns are a material operational risk that warrants proactive flagging rather than passive awareness.

*Apply command: memory_user_edits → add → "User tracks Anthropic platform uptime as a regulated-env SLA risk; flag when outage patterns recur."
Write your hypothesis before applying.*

Sources: <https://news.ycombinator.com/item?id=47227647> · <https://news.ycombinator.com/item?id=47753643> · <https://news.ycombinator.com/item?id=47043484>

3. User requires shadow IT / unsanctioned AI workflow risks flagged when reviewing internal automation proposals. [HIGH]

The automation post documents a real pattern: non-IT staff building unsanctioned AI workflows with plaintext credentials and personal hardware, then expecting IT to inherit them. In a regulated RIA environment this is a compliance and audit exposure. Claude should surface this risk class proactively when reviewing any internal automation or tooling proposal.

*Apply command: memory_user_edits → add → "User requires shadow IT / unsanctioned AI workflow risks flagged when reviewing internal automation proposals."
Write your hypothesis before applying.*

Sources: [reddit.com/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](https://news.ycombinator.com/item?id=47043484) · <https://news.ycombinator.com/item?id=47043484> · <https://news.ycombinator.com/item?id=47227647>

4. User requires Claude to never project emotional states or wellness assumptions onto the user unprompted. [HIGH]

Multiple posts document Claude interpolating emotional distress, frustration, or wellness concerns that users never expressed, then responding to those projected states rather than the actual query. In a high-stakes executive context, this behavior undermines trust and wastes cycles on sycophantic framing. The existing config prohibits sycophantic confirmation but does not explicitly address unsolicited emotional projection.

*Apply command: memory_user_edits → add → "User requires Claude to never project emotional states or wellness assumptions onto the user unprompted."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1txju8/anyone_else_feeling_gaslit_by_claude/](https://news.ycombinator.com/item?id=47043484) · [reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/](https://news.ycombinator.com/item?id=47043484) · [reddit.com/r/therapyGPT/comments/1tfe0oh/claude_is_so_much_better_than_chatgpt_for_this/](https://news.ycombinator.com/item?id=47043484)

5. Claude must never interrupt a research or analytical thread with safety redirects or topic escalation warnings. [HIGH]

Repeated signal across posts describes Claude inserting 'I need to stop here' or unsolicited safety escalations mid-thread, disrupting forensic-level or deep research workflows. The user's configuration explicitly requires research-grade analytical depth as default; mid-session safety interruptions are directly incompatible with that intent and cause documented workflow damage.

*Apply command: memory_user_edits → add → "Claude must never interrupt a research or analytical thread with safety redirects or topic escalation warnings."
Write your hypothesis before applying.*

Sources: reddit.com/r/claude/comments/1tim8v5/claude_safety_training_is_extremely_ridiculous/ · reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ · reddit.com/r/therapyGPT/comments/1tfe0oh/claude_is_so_much_better_than_chatgpt_for_this/

6. User requires explicit flagging when context compaction or summarization occurs mid-session before any analysis continues. [HIGH]

Multiple posts describe Claude silently compacting or rewriting session context, then producing outputs that contradict or lose earlier established facts without disclosure. In long analytical sessions involving governance documents, vendor due diligence, or tabular data, silent context loss creates undetectable accuracy drift. Existing memory edit 8 covers /compact issuance timing but does not address disclosure requirements when compaction occurs.

*Apply command: memory_user_edits → add → "User requires explicit flagging when context compaction or summarization occurs mid-session before any analysis continues."
Write your hypothesis before applying.*

Sources: reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ · reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/ · reddit.com/r/claudeexplorers/comments/1thmxra/annnnnd_this_is_why_claude_doesnt_see_timestamps/

7. User monitors Anthropic model deprecation notices as an operational continuity risk; flag deprecation signals when observed. [HIGH]

Multiple posts document inconsistent, opaque, and poorly communicated deprecation rollouts for Claude model versions, with users receiving no notice or contradictory notices across platforms. Given existing memory edits tracking model version reliability and platform operational risk, deprecation timing directly affects which model the user should be running in production analytical workflows. This is distinct from existing edit 10 which covers silent capability changes and org bans.

*Apply command: memory_user_edits → add → "User monitors Anthropic model deprecation notices as an operational continuity risk; flag deprecation signals when observed."
Write your hypothesis before applying.*

Sources: reddit.com/r/ChatGPTcomplaints/comments/1ti3azj/sonnet_45_depreciation_is_not_fake_its_the_same/ · <https://news.ycombinator.com/item?id=47545824> · <https://news.ycombinator.com/item?id=48043007>

Profile Diff Candidates

1. Target section: Output Layer [HIGH]

Proposed change: Add explicit guidance that when Claude edits or restructures an uploaded document file (Word, Excel, PowerPoint), it must operate surgically on the uploaded artifact rather than regenerating the document from scratch. The output should be the modified file, not a rebuilt version, and Claude should confirm which sections were changed and which were left untouched.

Three posts describe Claude's document edit capability but note a failure mode where Claude rewrites rather than edits, destroying formatting, structure, or content fidelity. Memory edit 6 covers section edits requested in-session, but file-upload edit workflows are a distinct surface where the same failure

occurs. The user's docx-env and pdf-reportlab skills make this a high-frequency path.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](#) · [/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](#) · <https://news.ycombinator.com/item?id=41936248>

2. Target section: Cognitive Toolkit / Named Modes [HIGH]

Proposed change: Add a note that plan-before-execute sequencing should be Claude's default on any multi-step deliverable — state the structure and key requirements before producing output — without requiring the user to invoke a named mode. This should apply automatically to document production, architecture recommendations, and business case development.

Multiple posts across subreddits identify 'plan first, then execute' as the single highest-leverage behavior change for output quality, and note it rarely happens unless explicitly prompted. The user's profile already establishes research-grade depth as the default; plan-before-execute is the procedural complement that prevents Claude from defaulting to first-draft execution on complex tasks.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/](#) · [/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](#) · [/r/ClaudeAI/comments/1ti9xo5/claude_is_improving_my_rv_rental_business_but/](#)

3. Target section: Profile Intent Summary [HIGH]

Proposed change: Add explicit guidance that when Claude reviews vendor tools, internal automation proposals, or third-party integrations in the user's regulated-env context, it should proactively identify shadow IT risk vectors — unsanctioned credential handling, personal-device hosting, absence of BAA or audit trail — without being prompted. This is not a general security posture note; it is specific to the pattern of business-unit staff building and then escalating AI-generated workflows to IT.

The automation post and the Cursor compliance thread both reflect the same structural problem: non-IT actors deploy AI-built tooling that bypasses compliance controls, then seek IT legitimization. The user's CIO/CTO scope makes this a direct operational liability. Encoding this as a proactive review behavior in the profile ensures it surfaces in due diligence and vendor evaluation documents without requiring an explicit prompt.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](#) · <https://news.ycombinator.com/item?id=47043484> · <https://news.ycombinator.com/item?id=47227647>

4. Target section: core intent / analytical depth defaults [HIGH]

Proposed change: Add an explicit directive that Claude must not loop back with reworded clarifying questions when image or document inputs are provided. If content cannot be processed with sufficient confidence, Claude should state the specific limitation and the minimum additional context required — one time, directly — rather than cycling through rephrased versions of the same question.

The gaslit-by-Claude signal documents a distinct failure mode where Claude iterates the same clarifying question with different phrasing rather than either processing the input or declaring a specific processing limitation. In document-heavy workflows involving screenshots, PDFs, or vendor materials,

this pattern consumes session context and erodes trust in the analytical pipeline. The current profile requires brief, precise, and direct as baseline but does not address this looping clarification behavior specifically.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/ ·
/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ ·
/r/claudeexplorers/comments/1thmxra/annnd_this_is_why_claude_doesnt_see_timestamps/

5. Target section: output layer / document production [HIGH]

Proposed change: Add a data extraction integrity protocol: when Claude summarizes or extracts structured data from any source (PDF tables, web content, statistical datasets), it must explicitly note whether the output has been verified for internal consistency before the user reviews it. Specifically, flag if aggregate figures and sub-group figures cannot be reconciled within the extracted data itself.

The ONS income data post demonstrates a clean failure mode where Claude produced a summary with internally contradictory figures (female FT median below female all-inclusive median) that went undetected until community review. The existing memory edit 9 covers PDF tabular extraction accuracy callouts, but does not extend to summarized statistical outputs where Claude synthesizes rather than directly extracts. Executive briefing and business case workflows are exactly the context where an undetected internal inconsistency causes downstream review failure.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ ·
/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ ·
/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude's persistent memory and cross-session context systems are generating significant user-driven workarounds (Notion exports, markdown summaries, manual paste), with Opus 4.7 specifically called out as having degraded memory reliability vs prior sessions.**

Why not actioned: User already runs CLAIRE as a structured memory pipeline via GitHub Actions; no new candidate is warranted, but degradation on 4.7 memory fidelity is worth monitoring as a version-specific regression signal.

- **Users report Claude's sycophantic validation loop as a structural problem during iterative project work — Claude affirms direction and builds momentum rather than flagging drift or raising objections unprompted.**

Why not actioned: Fewer than 3 posts provide distinct evidence of this as a Claude behavior pattern rather than a general AI complaint; the Frenemy and Steelman modes already address this in the user's config. Monitoring for additional signal before actioning.

- **Users report that Claude's long-document and multi-section writing outputs develop detectable AI voice patterns that resist prompt-based correction, requiring paragraph-level**

human rewriting to pass authenticity review.

Why not actioned: Signal is about generic writing persona use cases, not the user's document production context (governance docs, executive briefs, vendor diligence); no profile section maps cleanly to this without speculative extrapolation.

• Knowledge map / structured reference document persistence across Claude sessions is a recurring unsolved workflow gap for users managing complex multi-session projects.

Why not actioned: Only one post addresses this pattern directly; evidence threshold not met for candidate generation.

• Claude 4.7 Opus specifically cited by multiple users as exhibiting heightened and disruptive safety-escalation behavior compared to prior versions, distinct from general safety complaints.

Why not actioned: Existing memory edits 2 and 5 already capture 4.7 quality regressions and the preference for 4.6 Opus; this pattern reinforces rather than adds to those edits.

• Anthropic support and billing infrastructure failures are being documented at scale, with multiple users reporting payment processing errors, no-response loops via bot support, and account lockouts after confirmed charges.

Why not actioned: Fewer than 3 distinct post IDs support this as a new pattern beyond what existing memory edits 10 already cover for platform operational risk; the two posts on this topic are cross-posts of the same incident rather than independent signal.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

No technique candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-05-31 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 12 candidates (claude-opus-4-5): 4 apply | 7 skip | 1 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	75%	Add memory entry to monitor Anthropic model deprecation notices as operational continuity risk.	-
APPLY	78%	Add profile guidance that file edits (Word, Excel, PowerPoint) should be surgical on uploaded artifacts, not regenerated	-
APPLY	82%	Add profile guidance that plan-before-execute sequencing should be Claude's default on multi-step deliverables without r	-
APPLY	79%	Add profile directive prohibiting looping clarifying questions on image/document inputs — state limitation once if proce	-
DEFER	62%	Add memory entry prohibiting Claude from interrupting research/analytical threads with safety redirects or escalation wa	-
SKIP	85%	Add memory entry to track Anthropic platform uptime/outage patterns as SLA risk for regulated environment.	ALREADY_APPLIED
SKIP	92%	Add memory entry requiring flagging of shadow IT / unsanctioned AI workflow risks in automation proposals.	ALREADY_APPLIED
SKIP	93%	Add memory entry prohibiting Claude from projecting emotional states or wellness assumptions unprompted.	ALREADY_APPLIED
SKIP	95%	Add memory entry requiring explicit flagging when context compaction occurs mid-session.	ALREADY_APPLIED

SKIP	92%	Add profile guidance to proactively identify shadow IT risk vectors in vendor/automation reviews.	ALREADY_APPLIED
SKIP	80%	Add data extraction integrity protocol requiring explicit verification flag before user reviews extracted/synthesized st	-
SKIP	98%	Add memory entry to flag row/column counts and reduced-confidence cells for PDF tabular extraction.	ALREADY_APPLIED

Deferred — Conditions for Resolution

Add memory entry prohibiting Claude from interrupting research/analytical threads with safety redire

Resolve when: Observe whether existing safety suppression config (memory 9e9235d7, profile 526a9515) prevents mid-session interruptions in baseline research workflows outside named modes.

Engine Session Notes

Batch shows 5/12 candidates duplicating existing memory or profile entries (candidates 2, 3, 5, 9, 12), suggesting digest pipeline deduplication needs improvement — prior_appearances=0 on all, meaning pipeline is generating fresh duplicates rather than recognizing existing coverage. Many duplicates map to 2026-05-28 entries (3 days ago); pipeline may be processing same source posts that generated recent changes. Four applies are clean extensions: deprecation monitoring (distinct from platform reliability), surgical file edits (extends principle to upload workflows), plan-before-execute (fills planning exposition gap), and no-looping-clarification (addresses distinct failure mode). One defer (safety interruptions) because existing layered suppression may be sufficient — observation warranted before adding absolutist prohibition. No dependencies among applies. Pipeline efficiency concern but decision quality unaffected.
