

CLAIRE

Claude Configuration Signal Report — 2026-05-31

Signal Summary

Run date: 2026-05-31

Posts scanned: 259

Signal Breakdown

Category	Count
behavior_complaint	40
workflow_gap	38
feature_praise	94
competitor_gap	2
cross_platform	0
noise (dropped)	85

Developer Signal Disposition

138 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User uploads existing Word/Excel/PDF files for in-place editing; preserve source format in all outputs. [HIGH]

Multiple posts document users discovering late that Claude can edit and restructure uploaded documents while preserving file format — not just generate new ones. Given the user's established document production workflow (Word, PDF, PowerPoint, spreadsheet outputs with Waterton branding), explicitly flagging this as a standing expectation prevents Claude from defaulting to 'here is what I would change' prose instead of producing an edited file artifact.

*Apply command: memory_user_edits → add → "User uploads existing Word/Excel/PDF files for in-place editing; preserve source format in all outputs."
Write your hypothesis before applying.*

Sources: [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · [reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/) · <https://news.ycombinator.com/item?id=41936248>

2. User expects PDF/table data extraction to be flagged for verification; never treat extracted numerics as authoritative without caveat. [HIGH]

The PDF extraction thread explicitly warns that extraction is reliable for small tables but fails silently on larger documents — 19 correct extractions followed by an undetected error. For a user whose outputs enter formal review processes and regulatory environments, a silent numeric error in a financial statement or vendor contract is a high-severity failure. A standing memory edit surfaces the verification obligation automatically.

*Apply command: memory_user_edits → add → "User expects PDF/table data extraction to be flagged for verification; never treat extracted numerics as authoritative without caveat."
Write your hypothesis before applying.*

Sources: [reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/) · [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · <https://news.ycombinator.com/item?id=41936248>

3. User treats context degradation in long sessions as a reliability risk; flag when summarization may affect precision. [HIGH]

Multiple posts document Claude silently rewriting or compressing context in long sessions, causing factual drift and inverted summaries. For a user building governance documents and analytical timelines where precision is non-negotiable, silent context loss is a direct workflow threat, not a UX nuisance.

*Apply command: memory_user_edits → add → "User treats context degradation in long sessions as a reliability risk; flag when summarization may affect precision."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/) · [reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](https://www.reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/) · [reddit.com/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/](https://www.reddit.com/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/)

4. User requires Claude to avoid unsolicited emotional tone-reads or protective reframes when delivering direct analysis. [HIGH]

Multiple posts report Claude inserting unprompted emotional validation, wellness checks, or protective reframings that derail analytical exchanges. This user's configuration explicitly trades comfort for accuracy; sycophantic emotional interpolation is as counterproductive as hedge-everything responses.

*Apply command: memory_user_edits → add → "User requires Claude to avoid unsolicited emotional tone-reads or protective reframes when delivering direct analysis."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](https://www.reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/) · [reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/](https://www.reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/)

Profile Diff Candidates

1. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add explicit guidance that when the user uploads an existing document (Word, Excel, PDF, PowerPoint), the default behavior is in-place editing with format preservation — not commentary on what should change. Claude should produce the edited artifact directly unless the user specifies otherwise. For PDF source documents where data is being extracted into a spreadsheet, append a verification flag to the output noting any rows or cells where extraction confidence is reduced (dense tables, scanned pages, non-standard formatting).

The Reddit signal shows a consistent pattern of users — including technically sophisticated ones — discovering these capabilities only after months of manual workarounds. Since the user's document workflow is already codified in the profile, extending it to cover upload-and-edit as a first-class operation eliminates a likely friction point. The verification flag addition is specifically motivated by the financial statement and contract use cases where silent extraction errors carry regulatory exposure.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](#) · [/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](#) · <https://news.ycombinator.com/item?id=41936248>

2. Target section: Analytical Depth / Research-Grade Default [HIGH]

Proposed change: Add a note that when role framing or domain context is absent from a request, Claude should infer it from the established user profile rather than defaulting to generic framing. The user's CIO/CTO operating context, regulated environment, and document-review orientation are standing role context — they do not need to be restated per prompt. Responses should reflect that expertise level by default, not require explicit 'act as a senior technology executive' scaffolding.

The PromptEngineering thread on output inconsistency identifies 'no role = no expertise' as the primary cause of generic outputs, and multiple posts corroborate that explicit role framing produces materially better results. The user's profile already establishes research-grade depth as the default, but this diff tightens the connection between that intent and the specific failure mode — Claude reverting to generic register when a prompt arrives without role scaffolding, even though the profile already encodes who the user is.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/PromptEngineering/comments/1tgleu2/my_claude_outputs_were_inconsistent_for_weeks/](#) · [/r/PromptEngineering/comments/1tl9gk6/this_claude_prompt_gave_me_the-most-accurate/](#) · [/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](#)

3. Target section: Infrastructure / Tooling / Vendor Evaluation [HIGH]

Proposed change: Add a standing due diligence framework for AI tooling in regulated environments. When the user evaluates AI-adjacent vendors (coding assistants, API services, SaaS platforms), Claude should proactively surface compliance-relevant dimensions: BAA availability, FedRAMP/SOC2 status, data routing architecture, audit trail capabilities, and whether the vendor's cloud model is compatible with RIA/SEC data handling obligations — without being asked to run through a compliance checklist.

The Cursor compliance thread directly mirrors the user's operating context: regulated industry, remote workforce, and vendor tooling decisions that have compliance surface area. The user's profile already covers vendor due diligence as a primary use case, but the signal indicates AI tooling specifically generates a distinct category of compliance questions (HIPAA BAA, FedRAMP, audit trails, cloud routing) that differ from traditional software vendors. Adding this as a named sub-framework ensures Claude doesn't treat AI tool evaluation like generic SaaS procurement.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=47043484> · <https://news.ycombinator.com/item?id=44111298> · [/r/SaaS/comments/1tis9wd/claude_pro_vs_chatgpt_plus_for_vibe_coding_which/](#)

4. Target section: Core Analytical Defaults [HIGH]

Proposed change: Add an explicit instruction that when context compression or conversation summarization occurs, Claude must surface this proactively and flag any analytical claims that depend on earlier context that may have been compacted. Claude should not silently continue producing precision-dependent outputs after a context window event without acknowledging the reliability boundary.

The signal shows a documented pattern where Claude continues generating confident, precise outputs after context has been silently rewritten, which is uniquely dangerous in the user's work — governance documentation, vendor due diligence, and executive briefings cannot absorb unannounced factual drift. The user's current profile demands research-grade defaults; this adds the reliability disclosure layer that research-grade work requires.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](#) · [/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/](#)

5. Target section: Output and Document Production [HIGH]

Proposed change: Add a standing instruction that numerical summaries and data-derived tables must include an explicit source-fidelity check step: where Claude has synthesized or reformatted source data, it should note which figures are direct quotations from the source versus derived or rounded — particularly in briefing documents and executive one-pagers where a transposition error would survive downstream review.

Multiple posts document Claude producing data summaries where figures in the TL;DR contradict figures in the body of the same document, with the error going undetected until community review. The user's profile already requires audience translation without losing accuracy; this codifies a verification step at the output layer for quantitative content, which is directly relevant to the macro and financial analysis use cases.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](#) · [/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](#) · [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#)

6. Target section: Cognitive Toolkit / Mode Triggers [MEDIUM]

Proposed change: Add a clarifying note that the existing mode triggers — particularly Anchor and Focus — should also fire as self-checks when Claude detects it has asked the same or substantially similar clarifying question more than once in a session. Repeated rephrasing of unanswered questions is a documented failure mode that erodes trust in long analytical sessions; the mode system should absorb this rather than letting it accumulate silently.

The gaslit-by-Claude pattern in the signal describes Claude looping on the same clarifying question with different phrasing, which users experience as either evasion or processing failure. For a user running multi-hour governance or due-diligence sessions, this pattern is a direct productivity and trust cost. The existing Anchor and Focus modes are already designed to interrupt drift; extending them to catch repetitive question loops is a natural, non-structural addition.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](#) · [/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/](#) · [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#)

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Users report Claude becoming an enthusiasm accelerant — surfacing adjacent opportunities and expanding project scope beyond the original request, which increases workload rather than reducing it.**

Why not actioned: Only two posts directly describe this dynamic (RV rental business thread, home cinema anecdote in comments) — below the three-post evidence threshold.

- **Persistent memory architectures using Notion or external documents as a session-bridging layer are gaining traction as a workaround for context loss across long projects.**

Why not actioned: The signal comes primarily from one detailed guide post and associated comments — not yet corroborated across three independent posts — and the user's existing memory configuration may already address this need.

- **Claude usage cap and session continuity constraints are generating workflow disruption for users in long analytical or build sessions, with the 5-hour reset cited as a specific breakpoint.**

Why not actioned: This is a platform constraint, not a configuration issue — no memory edit or profile change can address session limits, and the user's profile intent does not include vibe coding or long uninterrupted build sessions as a primary use case.

- **Community discussion consistently flags LLM outputs as failing to sustain a distinctive analytical voice or argument structure across long-form documents, with AI-generated prose identifiable by its absence of argumentative logic.**

Why not actioned: Only two posts surface this pattern at the threshold level relevant to this user's document production use case; the third post (Pokémon/agent limitations) addresses autonomous task completion rather than prose quality. Does not meet the 3-post evidence threshold for the specific failure mode of analytical voice degradation in document outputs.

- **Claude increasingly interrupts research-oriented sessions with unsolicited protective stops — citing user wellbeing, potential harm, or escalating concern — even in clearly academic or professional contexts, and the frequency appears to have increased with recent model versions.**

Why not actioned: The user's profile already explicitly addresses this via the comfort-for-accuracy tradeoff and direct-voice defaults; a new candidate would duplicate existing configuration intent rather than add to it.

- **Shadow governance risk: non-technical staff are independently building Claude-assisted workflows with production data, plain-text credentials, and no IT review — then escalating to IT for infrastructure support after deployment.**

Why not actioned: This is an organizational governance pattern surfaced in the signal, not a Claude behavior pattern — it belongs in a policy or architecture review, not a configuration candidate.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

1. Plan-Before-Execute Prompting [MEDIUM]

Ask the AI to explicitly list what the output needs to accomplish or what criteria it should meet before producing the output. For example: 'Before writing, list the 4 things this email needs to do well. Then write it.' This forces a planning step that improves output quality by aligning the AI's goals with the user's intent before generation begins.

Test: Before your next writing or analysis request, prepend 'First list [N] things this response needs to do well, then produce it.' and compare the result to a direct request.

2. Examples Over Adjectives [MEDIUM]

Instead of describing desired tone or style with adjectives (e.g. 'professional,' 'friendly,' 'concise'), paste in concrete examples of text that has the quality you want. The AI infers the style from examples rather than interpreting vague descriptors, producing more accurate outputs.

Test: Replace style adjectives in your next prompt with one or two pasted examples of text you want the output to resemble, then compare the result to your adjective-based approach.

3. Pre-Task Context Elicitation [MEDIUM]

Before asking the AI to perform a task, first ask it what information it would need to do the task well. This surfaces missing context the user may not have thought to provide, reducing incomplete or generic outputs.

Test: On your next complex request, first ask 'What would you need to see or know to do this task well?' and provide what it requests before issuing the main task prompt.

4. Explicit Non-Invention Instruction [MEDIUM]

Include an explicit instruction telling the AI not to invent, assume, or fabricate information, and to ask for clarification when something is unknown. Phrased as 'don't assume anything and always ask for clarification rather than inventing details.' This is noted as effective across AI tools when accuracy is critical.

Test: Add 'Do not assume or invent any information; if you are uncertain about a detail, ask me for clarification instead' to your next high-stakes prompt and observe whether hallucinations or unsupported claims decrease.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-05-31 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 14 candidates (claude-opus-4-5): 3 apply | 9 skip | 2 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
----------	------------	-----------	------------

APPLY	72%	Establish that when user uploads existing documents (Word/Excel/PDF), Claude should preserve source format in outputs an	SOURCE_NEW
APPLY	74%	When role framing is absent, infer from established user profile rather than defaulting to generic framing.	SOURCE_NEW
APPLY	71%	Numerical summaries must include source-fidelity check distinguishing direct quotations from derived/rounded figures.	SOURCE_NEW
DEFER	58%	Profile change bundling in-place editing guidance with PDF extraction verification flags.	DEPENDENCY_UNRESOLVED, ALREADY_APPLIED
DEFER	52%	Anchor and Focus modes should fire as self-checks when Claude detects repeated clarifying questions.	HIGH_VARIANCE
SKIP	92%	Flag PDF/table data extraction for verification; never treat extracted numerics as authoritative without caveat.	ALREADY_APPLIED
SKIP	91%	Treat context degradation as a reliability risk; flag when summarization may affect precision.	ALREADY_APPLIED
SKIP	90%	Avoid unsolicited emotional tone-reads or protective reframes when delivering direct analysis.	ALREADY_APPLIED
SKIP	93%	Add standing due diligence framework for AI tooling with compliance dimensions (BAA, FedRAMP, data routing, audit trails	ALREADY_APPLIED
SKIP	91%	When context compression occurs, surface it proactively and flag claims depending on compacted context.	ALREADY_APPLIED
SKIP	78%	Plan-Before-Execute Prompting technique - ask AI to list what output needs to accomplish before producing it.	LOW_SIGNAL, SOURCE_NEW
SKIP	78%	Examples Over Adjectives technique - use concrete examples instead of style adjectives.	LOW_SIGNAL, SOURCE_NEW
SKIP	80%	Pre-Task Context Elicitation technique - ask AI what information it needs before performing task.	LOW_SIGNAL, SOURCE_NEW
SKIP	85%	Explicit Non-Invention Instruction technique - tell AI not to invent and to ask for clarification.	LOW_SIGNAL, ALREADY_APPLIED, SOURCE_NEW

Deferred — Conditions for Resolution

Profile change bundling in-place editing guidance with PDF extraction verification flags.

Resolve when: Evaluate outcome of 6adf96e4 (in-place editing memory edit). If that change proves insufficient, unbundle this candidate's editing guidance from its verification component and re-evaluate.

Anchor and Focus modes should fire as self-checks when Claude detects repeated clarifying questions.

Resolve when: Awaiting user confirmation that the repetitive clarifying question pattern is occurring in their sessions. If confirmed, reframe candidate with specific trigger criteria before applying.

Engine Session Notes

This batch exhibits a significant deduplication issue: 5 of 14 candidates (36%) are functionally duplicative of changes already in memory state. The digest pipeline appears to be proposing candidates without adequate filtering against current configuration. Recommend reviewing the assembler's memory consistency check - the meta note about missing 'summary' fields in change_log may be contributing to this gap. The technique track (4 candidates) uniformly failed signal threshold (all at 0.57 strength, single sources, MEDIUM confidence). Consider raising the minimum signal threshold for technique candidates or requiring multi-source convergence before surfacing. Three candidates are applied: (1) in-place document editing format preservation, (2) role inference from profile context, (3) source-fidelity tracking for numerical outputs. All three are coherent extensions of existing configuration with low regression risk. Two candidates deferred: (1) bundled profile change pending resolution of overlapping memory edit, (2) mode trigger extension awaiting user confirmation of the underlying pattern. Both deferrals include specific resolution conditions. No contradictions or regressions detected. The pipeline is producing coherent recommendations but at reduced efficiency due to the deduplication gap. No escalation required - confidence distribution is appropriately bimodal, not flat.
