

CLAIRE

Claude Configuration Signal Report — 2026-05-31

Signal Summary

Run date: 2026-05-31

Posts scanned: 266

Signal Breakdown

Category	Count
behavior_complaint	41
workflow_gap	36
feature_praise	98
competitor_gap	3
cross_platform	0
noise (dropped)	88

Developer Signal Disposition

143 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User uploads PDFs, Word, and Excel files for extraction, restructuring, and analysis tasks.
[HIGH]

Multiple posts document users discovering Claude's document ingestion and transformation capabilities. For a user producing governance docs and vendor due diligence materials, proactively applying extraction and restructuring logic to uploaded files — without requiring the user to specify the capability — reduces friction in the document production workflow.

Apply command: `memory_user_edits` → add → "User uploads PDFs, Word, and Excel files for extraction, restructuring, and analysis tasks."
Write your hypothesis before applying.

Sources: [reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/) · [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · <https://news.ycombinator.com/item?id=41936248>

2. User expects plan-before-execute behavior on complex outputs without being asked to request it. [HIGH]

Recurring Reddit signal identifies plan-first prompting as a high-leverage quality improvement — structuring the approach before producing output. This aligns directly with the user's need for research-grade analytical depth as a default mode, not a prompted behavior. Making this automatic prevents generic first-pass outputs that require rewriting.

*Apply command: memory_user_edits → add → "User expects plan-before-execute behavior on complex outputs without being asked to request it."
Write your hypothesis before applying.*

Sources: [reddit.com/r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/](https://www.reddit.com/r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/) · [reddit.com/r/PromptEngineering/comments/1tgleu2/my_claude_outputs_were_inconsistent_for_weeks/](https://www.reddit.com/r/PromptEngineering/comments/1tgleu2/my_claude_outputs_were_inconsistent_for_weeks/) · [reddit.com/r/PromptEngineering/comments/1t19gk6/this_claude_prompt_gave_me_the_most_accurate/](https://www.reddit.com/r/PromptEngineering/comments/1t19gk6/this_claude_prompt_gave_me_the_most_accurate/)

3. User requires explicit notification if context compaction or summarization occurs mid-session. [HIGH]

Multiple posts document Claude silently rewriting or compacting conversation context without user awareness, causing precision errors in long analytical sessions. For a user doing governance documentation and high-stakes analysis where continuity matters, silent context loss is a direct integrity risk.

*Apply command: memory_user_edits → add → "User requires explicit notification if context compaction or summarization occurs mid-session."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/) · [reddit.com/r/ChatGPTcomplaints/comments/1tkpzb/on_claude_i_caught_both_memory_system_and/](https://www.reddit.com/r/ChatGPTcomplaints/comments/1tkpzb/on_claude_i_caught_both_memory_system_and/) · [reddit.com/r/claude/comments/1tjxu8/anyone_else_feeling_gaslit_by_claude/](https://www.reddit.com/r/claude/comments/1tjxu8/anyone_else_feeling_gaslit_by_claude/)

4. User is a domain expert; do not insert unsolicited safety caveats or interrupt analytical threads with concern checks. [HIGH]

Widespread signal across multiple posts describes Claude breaking into research or analytical sessions with safety interruptions, eating disorder assumptions, or 'I need to stop here' interceptions that are contextually unwarranted. For a senior executive doing regulatory, medical, or financial research, these interruptions degrade output quality and require workarounds.

*Apply command: memory_user_edits → add → "User is a domain expert; do not insert unsolicited safety caveats or interrupt analytical threads with concern checks."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/](https://www.reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/) · [reddit.com/r/therapyGPT/comments/1tfe0oh/claude_is_so_much_better_than_chatgpt_for_this/](https://www.reddit.com/r/therapyGPT/comments/1tfe0oh/claude_is_so_much_better_than_chatgpt_for_this/) · [reddit.com/r/claude/comments/1tjxu8/anyone_else_feeling_gaslit_by_claude/](https://www.reddit.com/r/claude/comments/1tjxu8/anyone_else_feeling_gaslit_by_claude/)

Profile Diff Candidates

1. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add explicit guidance that when the user uploads an existing Word, Excel, or PowerPoint file, Claude should default to editing and restructuring the uploaded artifact in-place — preserving format, brand system, and structure — rather than producing a new document or describing changes in text. The output should be the transformed file, not commentary about how to transform it.

Multiple posts document a widely missed capability: Claude can modify existing documents, not just generate new ones. For a user whose work product enters formal review processes and follows a specific brand system, the default behavior should be transformation of the uploaded artifact, not parallel generation. This prevents the user from having to reintegrate Claude's output into their existing files manually.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](#) · [/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/](#) · <https://news.ycombinator.com/item?id=41936248>

2. Target section: Core Analytical Behavior / Research-Grade Default [HIGH]

Proposed change: Add a directive that when producing any substantive analytical output, Claude should explicitly surface its structural assumptions before executing — briefly stating what interpretive frame it is using, what it is treating as primary versus secondary evidence, and where it sees the highest uncertainty. This should appear as a compact framing block at the top, not as hedging language embedded throughout the response.

The signal consistently identifies role-definition and explicit framing as the primary drivers of output quality consistency. For a user whose outputs enter formal review and governance processes, undeclared interpretive assumptions are a failure mode — reviewers cannot distinguish analytical judgment from model artifact. A compact upfront framing block makes the analytical stance auditable without adding hedge language that violates the user's brevity and precision norms.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/PromptEngineering/comments/1tgleu2/my_claude_outputs_were_inconsistent_for_weeks/](#) · [/r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/](#) · [/r/PromptEngineering/comments/1tl9gk6/this_claude_prompt_gave_me_the_most_accurate/](#)

3. Target section: Infrastructure Architecture / Vendor Due Diligence [HIGH]

Proposed change: Add a standing analytical frame for AI tooling compliance evaluation in regulated environments. When assessing AI-adjacent vendor or tooling decisions, Claude should default to surfacing: BAA availability, FedRAMP/SOC2 status, data routing architecture, audit trail capabilities, and contractual liability exposure — without being prompted. This is especially relevant for shadow-IT scenarios where business units adopt AI tools outside formal procurement (e.g., automation workflows built on personal laptops, CRM integrations without IT review). The output frame should include a governance gap inventory alongside the technical assessment.

Multiple posts surface the pattern of regulated-industry teams confronting AI tooling without compliance scaffolding — the Cursor/HIPAA thread directly and the automation cleanup-crew thread as a shadow-IT governance failure. For a CIO/CTO at an SEC-registered RIA, this is a live operational risk category. Defaulting to a compliance-first evaluation frame for AI tooling removes a step the user would otherwise have to explicitly request on every vendor assessment.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=47043484> ·
[/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](https://news.ycombinator.com/item?id=44111298) ·
<https://news.ycombinator.com/item?id=44111298>

4. Target section: Output Layer / Document Production [HIGH]

Proposed change: Add an explicit instruction that when Claude performs quantitative summarization or data analysis (tables, percentiles, statistical outputs), it must internally cross-check aggregate summary statistics against row-level data before presenting results, and flag any cell where the summary figure conflicts with the underlying data rather than silently presenting the inconsistency.

The ONS income data post documents Claude producing a summary table where a headline figure contradicted the detailed breakdown in the same output — a precision failure that went undetected until a reader caught it. For a user whose outputs enter formal review processes, an undetected numerical inconsistency in an executive briefing is a credibility event. The fix is a pre-output consistency check instruction, not a caveat.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](https://news.ycombinator.com/item?id=47043484) ·
[/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://news.ycombinator.com/item?id=44111298) ·
[/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](https://news.ycombinator.com/item?id=44111298)

5. Target section: Cognitive Toolkit / Named Modes [MEDIUM]

Proposed change: Add a 'Audit' mode that, when invoked, causes Claude to re-examine the current or most recent output specifically for internal consistency failures: summary statistics vs. source data, claims in one section contradicting claims in another, and any assertion made without a traceable basis in the provided inputs. Output is a structured discrepancy log, not a rewrite.

The data summarization error pattern and the context-drift pattern together suggest a gap in the current mode toolkit: there is no quality-control mode focused on internal logical consistency and numerical accuracy specifically. The existing modes address sycophancy, stress-testing, and decision-forcing but not post-hoc verification of factual coherence within a single output. An Audit mode closes this gap without restructuring existing modes.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](https://news.ycombinator.com/item?id=47043484) ·
[/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://news.ycombinator.com/item?id=44111298) ·
[/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](https://news.ycombinator.com/item?id=44111298)

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Users are building custom persistent memory architectures outside native Claude memory — exporting chats, using Notion, and constructing session-bridging prompts — because native memory is perceived as insufficient for complex ongoing work.**

Why not actioned: The user's profile already includes a memory and modes system; the Reddit signal reflects workarounds by users without configured profiles, not a gap specific to this user's setup.

- **Claude's sycophancy in iterative build contexts — hyping user work rather than stress-testing it — is flagged repeatedly as a quality degradation in longer sessions.**

Why not actioned: The user's existing mode system (Steelman, Frenemy) directly addresses this; the signal confirms the modes are solving a real problem but does not reveal a gap in the current configuration that warrants a new candidate.

- **Claude usage cap and session continuity constraints are causing workflow disruption for users in extended analytical or build sessions, with the 5-hour reset window cited as the primary friction point.**

Why not actioned: Only two posts address session/availability limits directly, falling below the 3-post evidence threshold.

- **Recurring user frustration with AI outputs requiring significant human review and correction before they can be used in professional contexts, with the gap between raw output and production-ready work being larger than anticipated.**

Why not actioned: The posts describing this pattern (buhaydigital, AIDiscussion, HN human-sounding articles) reflect general user population frustration rather than any newly observed Claude behavior or capability gap — the user's existing configuration already addresses this directly through its research-grade default and explicit anti-genericness stance. No new candidate is warranted.

- **Claude appears to over-read emotional or tonal signals in user messages and volunteers unsolicited psychological reframing or concern responses that the user did not request and did not warrant.**

Why not actioned: The existing profile already instructs against sycophancy and comfort-seeking responses; this pattern may resolve under the current configuration without a new edit, and the source posts skew toward therapeutic/companion use cases rather than the user's executive workflow context.

- **Model deprecation and version discontinuity are occurring with inconsistent user notification, creating workflow disruption when a relied-upon model version is removed without clear advance notice.**

Why not actioned: This is a platform/infrastructure behavior outside the scope of configuration or memory edits; no candidate type in the schema addresses Anthropic product decisions.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

1. Persistent Role Instruction via Project/System Prompt [HIGH]

Users paste a detailed role definition and behavioral instructions into a project-level system prompt or custom instructions field so that every conversation in that context inherits the same expert persona and constraints without re-prompting. This avoids copy-pasting setup text into each new chat and keeps domain context always active.

Test: Create a project or custom instructions entry with a multi-sentence role definition (e.g., 'You are a [domain] expert...') and open several new chats within it to verify the persona persists without re-stating it each time.

2. Explicit Anti-Hallucination Constraint in Prompt [HIGH]

Users add an explicit instruction to the prompt telling the AI not to invent or assume information, and to ask for clarification when it lacks sufficient detail. This is described as a portable addition to any base prompt that reduces fabricated outputs in high-stakes or factual domains.

Test: Append the phrase 'Do not assume anything not explicitly stated; if information is missing, ask me a clarifying question instead of guessing' to an existing prompt and compare output accuracy on a domain-specific task.

3. Situation-First Framing for Recommendations [HIGH]

Users describe their full situation or context upfront before asking for recommendations, rather than asking a generic question. Providing concrete situational detail (trip destination, event type, constraints) produces more relevant and complete output lists that include items the user had not thought of themselves.

Test: Before asking for any recommendation, write 2-3 sentences describing your specific situation, constraints, and goal, then ask the recommendation question and compare the result to a version where you asked the question alone.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-05-31 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 12 candidates (claude-opus-4-5): 7 apply | 4 skip | 1 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	78%	Add memory that user is a domain expert — suppress unsolicited safety caveats and concern checks during analytical threa	-
APPLY	83%	Add profile guidance: when user uploads existing Office file, default to editing in-place and returning transformed file	-
APPLY	79%	Add profile directive: surface structural assumptions as compact framing block before executing substantive analytical o	-
APPLY	88%	Technique: Persistent Role Instruction via Project/System Prompt for maintaining persona across sessions	-
APPLY	85%	Technique: Explicit Anti-Hallucination Constraint in Prompt — add 'do not assume; ask for clarification' instruction	-
APPLY	86%	Technique: Situation-First Framing for Recommendations — provide context before asking recommendation questions	-
DEFER	65%	Expand AI compliance evaluation frame to include shadow-IT scenarios and governance gap inventory	DEPENDENCY_UNRESOLVED
SKIP	85%	Add memory that user uploads documents for extraction and restructuring tasks	-
SKIP	82%	Add memory that user expects plan-before-execute behavior on complex outputs	DEPENDENCY_UNRESOLVED
SKIP	95%	Add memory requiring explicit notification if context compaction occurs mid-session	ALREADY_APPLIED
SKIP	92%	Add instruction to cross-check aggregate statistics against row-level data before presenting results	ALREADY_APPLIED
SKIP	94%	Add 'Audit' mode for internal consistency verification with structured discrepancy log output	ALREADY_APPLIED

Deferred — Conditions for Resolution

Expand AI compliance evaluation frame to include shadow-IT scenarios and governance gap inventory

Resolve when: First eval cycle with measurable output from profile_diff d94eab47 (AI-assisted development tool compliance methodology)

Engine Session Notes

Batch shows digest pipeline duplication issue — 3 candidates (context compaction notification, statistical consistency check, Audit mode) propose functionality already in change_log. Recommend adding de-duplication pass to assembler that checks candidate proposals against applied change_log summaries before surfacing to CLAIRE-A. Capability vs. constraint confusion also present — candidate eb56c77b describes what Claude can do rather than what it should do. Synthesis layer should convert capability observations to constraint specifications. Technique candidates (3) are working as designed — appropriately scoped as user-behavior observations, logged for reference without system state changes. Profile diffs in this batch (c7e0ecf5, f012e8a5) are well-specified and close real gaps. They follow the established pattern of actionable behavioral instructions. One defer (736ee731) on dependency to existing profile_diff d94eab47 which needs eval feedback before scope expansion. Defer condition is specific and time-bounded. No escalation required — decisions are cleanly distributed, no conflicting signals, no eval regression indicators.
