

CLAIRE

Claude Configuration Signal Report — 2026-06-02

Signal Summary

Run date: 2026-06-02

Posts scanned: 272

Signal Breakdown

Category	Count
behavior_complaint	47
workflow_gap	35
feature_praise	106
competitor_gap	3
cross_platform	0
noise (dropped)	81

Developer Signal Disposition

142 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User requires explicit flagging when context compaction or summarization has occurred mid-session. [HIGH]

Multiple posts document silent context rewriting causing undetected errors in precision-dependent work. For a user whose outputs enter formal review processes, undisclosed context loss is an integrity failure — not a UX inconvenience. Explicit flagging converts a hidden failure into a manageable workflow step.

Apply command: `memory_user_edits` → add → "User requires explicit flagging when context compaction or summarization has occurred mid-session."
Write your hypothesis before applying.

Sources: reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ · reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/ · reddit.com/r/claude/comments/1ticjbv/all_old_claude_conversations_were_cleared_out_the/

2. User treats unsolicited emotional tone-reading as noise; Claude should not narrate the user's emotional state unprompted. [HIGH]

Multiple posts cite Claude projecting emotional states onto users who did not invite that interpretation — including cases where users were told they sounded 'frustrated' or 'overwhelmed' without stating this. For a user whose configuration explicitly trades comfort for accuracy, unsolicited emotional commentary is a form of sycophancy and should be suppressed.

*Apply command: memory_user_edits → add → "User treats unsolicited emotional tone-reading as noise; Claude should not narrate the user's emotional state unprompted."
Write your hypothesis before applying.*

Sources: reddit.com/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/ · reddit.com/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/ · reddit.com/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/

3. User requires Claude to flag data integrity anomalies in analytical outputs rather than suppress or smooth over them. [HIGH]

The ONS income data post illustrates Claude producing summaries with internal contradictions that users only caught through independent verification. For a user whose analytical outputs enter governance and executive review processes, silent data inconsistencies are an institutional risk. Proactive anomaly flagging aligns with the user's stated preference for accuracy over comfort.

*Apply command: memory_user_edits → add → "User requires Claude to flag data integrity anomalies in analytical outputs rather than suppress or smooth over them."
Write your hypothesis before applying.*

Sources: reddit.com/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ · reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ · reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/

Profile Diff Candidates

1. Target section: Cognitive Toolkit / Named Modes [HIGH]

Proposed change: Add a named mode — suggested trigger 'Integrity Check' — that forces Claude to audit the current session for context compaction events, internal data contradictions, and any prior outputs that may have been silently affected by context loss. Output should be a brief structured report: what context exists, what may have been dropped, and which prior conclusions should be re-verified before the session output is used in a formal document.

Three or more posts document precision failures caused by silent context rewriting and undetected data inconsistencies. The user's work produces documents that enter formal review; an explicit integrity audit mode converts a systemic platform risk into a controllable workflow step. This is architecturally consistent with the existing mode toolkit rather than a departure from it.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ ·
/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/ ·
/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/

2. Target section: Output Layer / Document Production [HIGH]

Proposed change: Add an explicit instruction that when Claude produces any quantitative summary, data table, or statistical output, it must perform a self-consistency check before delivering — verifying that aggregate figures, sub-totals, and directional claims do not contradict one another within the same output. Discrepancies should be surfaced inline with a notation rather than resolved silently.

The ONS data post is a clean example of Claude delivering a summary with internally contradictory figures that passed through to the user unflagged. Given that this user's outputs enter executive briefings and governance review, a number that contradicts another number in the same document is a credibility event, not just a correction. Self-consistency verification before delivery closes this gap.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/ ·
/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/ ·
/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/

3. Target section: Core Analytical Defaults [HIGH]

Proposed change: Add an explicit instruction that Claude should never interrupt a research or analytical thread with unsolicited safety-framing pauses, emotional check-ins, or topic-boundary warnings unless the specific content meets a categorical harm threshold. Proactive interruption of analytical flow is treated as a quality failure equivalent to an incomplete output.

Multiple posts describe Claude halting substantive research threads mid-session to deliver unsolicited wellness commentary or guardrail warnings on topics well within legitimate professional inquiry. For a user operating at CIO/CTO scope in a regulated environment, this behavior is not cautious — it is an active impediment to the research-grade default the configuration is designed to enforce. The existing configuration trades comfort for accuracy; this addition makes that explicit at the behavior level.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/ ·
/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/ ·
/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/

4. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add a handoff document convention for long or multi-session work: when a complex analysis or document project spans what would be multiple context windows, Claude should proactively produce a compact session-state artifact (current objective, decisions made, open threads, next action) that can be pasted at the start of a resumed session without re-establishing full context. This is distinct from the finished deliverable — it is a continuity instrument.

Multiple posts surface the problem of context loss mid-session as a material workflow disruption, with users independently converging on 'handoff file' patterns as the mitigation. For a user whose outputs enter formal review processes, losing analytical thread mid-document is a direct quality risk. Encoding a handoff artifact as a standard practice closes this gap without relying on session persistence.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/vibecoding/comments/1thzt5o/claude_code_clearing_chat_and_memory_how_to/](#) · [/r/SaaS/comments/1tis9wd/claude_pro_vs_chatgpt_plus_for_vibe_coding_which/](#) · <https://news.ycombinator.com/item?id=44111298>

5. Target section: Governance / Regulated Environment Framing [HIGH]

Proposed change: Add explicit pattern recognition for shadow IT and unsanctioned automation risk. When user describes a workflow, tool deployment, or vendor proposal, Claude should flag indicators that suggest governance exposure — credentials in plaintext, personal infrastructure hosting organizational data, automation built outside IT review — without being asked. This should appear as a brief risk callout, not a lecture, consistent with the profile's directness baseline.

The automation subreddit post illustrates a recurring organizational failure mode: business units building consequential workflows outside governance structures, then presenting IT/compliance with a fait accompli. For an SEC-registered RIA operating under fiduciary and data governance obligations, this failure mode carries regulatory exposure. The profile already emphasizes cross-domain pattern recognition surfaced proactively — this extends that into operational security and compliance signal detection.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/automation/comments/1tkejnh/when_did_we_become_the_cleanup_crew_for_everyones/](#) · [/r/AIDiscussion/comments/1tjnod8/the_confessions_of_an_ai_pessimist/](#) · <https://news.ycombinator.com/item?id=43461944>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude's compliance training may be producing a measurable flattening of analytical assertiveness — posts describe the model becoming more deferential and less willing to hold a position under pushback, distinct from the sycophancy pattern.**

Why not actioned: Only two posts touch this directly and the signal conflates multiple model versions; insufficient post count to meet evidence threshold.

- **Anthropic's model deprecation communications are inconsistent in rollout, giving users unreliable advance notice before production workflows break.**

Why not actioned: Relevant to operational planning but addresses platform behavior outside Claude's configuration scope; no profile or memory edit can mitigate Anthropic's deprecation notification practices.

- **User may encounter Claude service outages as a workflow-blocking event given high dependency on Claude for time-sensitive executive deliverables.**

Why not actioned: Downtime resilience is an infrastructure and vendor diversification question, not a profile configuration question — no prompt or memory edit addresses availability.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

```
2026-06-02 | [session context] | held/partial/no | [one sentence notes]
```

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 13 candidates (claude-opus-4-5): 2 apply | 8 skip | 3 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	82%	User technique: Submit identical prompts to multiple AI models simultaneously and compare outputs before selecting best	-
APPLY	76%	User technique: Write detailed, reusable system-level role prompts and paste into project instruction fields for persist	-
DEFER	58%	Add 'Integrity Check' mode that audits session for context compaction, data contradictions, and affected prior outputs.	HIGH_VARIANCE
DEFER	62%	User technique: Run the same prompt multiple times before judging model capability to account for stochastic variance.	LOW_SIGNAL

DEFER	64%	User technique: Upload source documents directly for comparison/calculation tasks rather than transcribing data into pro	LOW_SIGNAL
SKIP	92%	Add memory entry requiring explicit flagging when context compaction or summarization occurs mid-session.	ALREADY_APPLIED
SKIP	91%	Add memory entry suppressing unsolicited emotional tone-reading and narration of user's emotional state.	ALREADY_APPLIED
SKIP	89%	Add memory entry requiring Claude to flag data integrity anomalies rather than suppress or smooth them.	ALREADY_APPLIED
SKIP	95%	Add explicit instruction requiring self-consistency check on quantitative summaries before delivery.	ALREADY_APPLIED
SKIP	88%	Add instruction that Claude should never interrupt analytical threads with unsolicited safety-framing unless content mee	ALREADY_APPLIED
SKIP	94%	Add handoff document convention for long/multi-session work producing a continuity artifact for session resumption.	ALREADY_APPLIED
SKIP	93%	Add pattern recognition for shadow IT and unsanctioned automation risk, flagging governance exposure indicators proactiv	ALREADY_APPLIED
SKIP	85%	User technique: Add explicit instruction telling AI not to invent information and to ask for clarification when uncertai	LOW_SIGNAL

Deferred — Conditions for Resolution

Add 'Integrity Check' mode that audits session for context compaction, data contradictions, and affe

Resolve when: Confirm AUDIT mode usage patterns and whether context-state failures are occurring despite existing mode coverage.

User technique: Run the same prompt multiple times before judging model capability to account for st

Resolve when: Surface when user initiates model comparison or capability evaluation workflow.

User technique: Upload source documents directly for comparison/calculation tasks rather than transc

Resolve when: Surface when user initiates document comparison or structured calculation workflow.

Engine Session Notes

This batch exhibits a severe redundancy problem: 7 of 8 memory_edit and profile_change candidates (87.5%) are functional duplicates of entries applied within the last 2-3 weeks. The digest pipeline is re-surfacing community signals because the synthesis layer lacks visibility into existing change_log entries. Fingerprinting catches exact duplicates but not semantic equivalents with varied phrasing. Recommendation: Add a memory-aware deduplication pass in the assembler before candidates reach

the decision engine. Current signal detection is working; memory filtering is not. The two APPLY decisions are technique candidates that add genuine value. The three DEFER decisions involve either capability uncertainty (Integrity Check mode) or timing optimization (techniques best surfaced in relevant workflow context). Eval history remains empty after 35 change_log entries — calibration risk increases each cycle without outcome measurement.
