

CLAIRE

Claude Configuration Signal Report — 2026-06-02

Signal Summary

Run date: 2026-06-02

Posts scanned: 133

Signal Breakdown

Category	Count
behavior_complaint	20
workflow_gap	13
feature_praise	49
competitor_gap	0
cross_platform	1
noise (dropped)	50

Developer Signal Disposition

75 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

No memory edit candidates this cycle. Signal below evidence threshold or no patterns identified.

Profile Diff Candidates

1. Target section: Vendor Due Diligence / Governance Documentation [MEDIUM]

Proposed change: Add a reusable analytical framework for AI tooling compliance evaluations in regulated environments. When the user is assessing AI vendors (coding assistants, API providers, SaaS tools), Claude should default to structuring the analysis around: (1) certification status — HIPAA BAA, SOC 2 Type II, FedRAMP, SEC-relevant controls; (2) data residency and routing architecture; (3) audit trail availability and export format; (4) contractual risk exposure if certification is absent; and (5) compensating controls that peer organizations in regulated industries have deployed when the vendor

is non-compliant. This should be surfaced as a structured comparison, not a narrative hedge.

The Cursor compliance post reflects a recurring evaluation pattern directly relevant to an SEC-registered RIA with a remote model — AI tooling procurement decisions where the compliance gap is predictable and the analytical structure should not have to be rebuilt from scratch each time. The user's profile already lists vendor due diligence as a core use case; this sharpens the default framework for the AI-tooling-in-regulated-environments sub-case specifically. Without this, Claude may produce generic vendor comparison output rather than the compliance-first framing this environment requires.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=47043484>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Users report that Claude produces detectable AI writing patterns even when explicitly prompted for human-sounding output, with paragraph-level stylistic uniformity being the primary tell.**

Why not actioned: Only one post supports this pattern, and the user's profile use cases (governance docs, executive briefings, analytical output) are not attempting to pass as human writing — they are attributed institutional documents where AI-assistance is not being concealed, so the failure mode described does not map to a real gap in this configuration.

- **Claude demonstrates persistent failures in long-horizon autonomous task completion requiring stateful planning across many sequential decisions.**

Why not actioned: Only one post supports this pattern, and while agentic use cases are adjacent to this user's workflow, the evidence threshold of 3 posts is not met.

- **Claude Pro usage limits and throttling are a persistent operational friction point for heavy users, with multiple posts describing hit walls mid-session, faster-than-advertised limit consumption during peak hours, and outright downtime events lasting over an hour.**

Why not actioned: This is an infrastructure/service reliability complaint with no signal specific to the user's configuration, analytical needs, or prompt behavior — no profile or memory change would mitigate a platform-level capacity constraint.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-06-02 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 1 candidates (claude-opus-4-5): 0 apply | 1 skip | 0 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
SKIP	82%	Add a structured 5-point compliance evaluation framework for AI vendor assessments covering certifications, data residen	LOW_SIGNAL, SOURCE_NEW, ALREADY_APPLIED

Engine Session Notes

Single-candidate batch. Candidate correctly identifies a relevant domain (AI vendor compliance evaluation) but substantially duplicates three existing memory entries applied within the last 10 days. The digest pipeline may benefit from stronger deduplication against recent change_log entries - assembler notes flag missing 'summary' fields on memory_edit entries which may contribute to weak matching. The compensating controls element is the only novel signal but lacks reinforcement; if this represents a real gap it should resurface with stronger evidence. No escalation required.
