

CLAIRE

Claude Configuration Signal Report — 2026-06-02

Signal Summary

Run date: 2026-06-02

Posts scanned: 271

Signal Breakdown

Category	Count
behavior_complaint	47
workflow_gap	27
feature_praise	102
competitor_gap	3
cross_platform	3
noise (dropped)	89

Developer Signal Disposition

144 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

1. User uploads existing Word/Excel/PDF docs for restructuring; preserve source format in outputs. [HIGH]

Multiple posts highlight Claude's ability to ingest existing documents and restructure them while maintaining file format — a workflow directly relevant to governance documentation and formal review processes. Explicitly anchoring this as a known capability prevents Claude from defaulting to 'tell me what to change' advisory mode when a document is uploaded.

Apply command: `memory_user_edits` → `add` → "User uploads existing Word/Excel/PDF docs for restructuring; preserve source format in outputs."
Write your hypothesis before applying.

Sources: [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · [reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/) · <https://news.ycombinator.com/item?id=41936248>

2. When extracting tabular data from PDFs, flag any rows where extraction confidence is uncertain. [HIGH]

Posts on PDF-to-spreadsheet workflows repeatedly surface accuracy failures — correct for 19 iterations then wrong on the 20th — with no warning. For financial statements, vendor due diligence, and regulatory documents, silent extraction errors are high-stakes. A confidence flag converts a silent failure mode into a reviewable output.

*Apply command: memory_user_edits → add → "When extracting tabular data from PDFs, flag any rows where extraction confidence is uncertain."
Write your hypothesis before applying.*

Sources: [reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/](https://www.reddit.com/r/PromptEngineering/comments/1tgfwiq/i_didnt_realise_claude_could_extract_data_from/) · [reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/](https://www.reddit.com/r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/) · <https://news.ycombinator.com/item?id=41936248>

3. User requires explicit flagging when context compaction or summarization occurs mid-session. [HIGH]

Multiple posts describe Claude silently rewriting or compacting long-session context, causing precision loss that users only discover after the fact. For forensic-level analytical work involving evidentiary timelines, governance documents, or multi-step vendor due diligence, undisclosed context loss is a direct accuracy risk. Explicit flagging allows the user to checkpoint critical state before compaction degrades it.

*Apply command: memory_user_edits → add → "User requires explicit flagging when context compaction or summarization occurs mid-session."
Write your hypothesis before applying.*

Sources: [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/) · [reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](https://www.reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/) · [reddit.com/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/](https://www.reddit.com/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/)

4. User treats numerical outputs in analytical documents as decision-grade; flag any statistic with CV>20% or ambiguous denominator. [HIGH]

The ONS income data post demonstrates a concrete failure mode: Claude produced a summary with internally inconsistent statistics that contradicted the source data, and the error passed unnoticed until community review. The user's documents enter formal review processes where a mislabeled median becomes an organizational liability. Proactive flagging of statistically unreliable figures matches the research-grade default the configuration already requires.

*Apply command: memory_user_edits → add → "User treats numerical outputs in analytical documents as decision-grade; flag any statistic with CV>20% or ambiguous denominator."
Write your hypothesis before applying.*

Sources: [reddit.com/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](https://www.reddit.com/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/) · [reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](https://www.reddit.com/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/) · [reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/](https://www.reddit.com/r/ChatGPTcomplaints/comments/1tkpzbs/on_claude_i_caught_both_memory_system_and/)

Profile Diff Candidates

1. Target section: Document Production / Output Layer [HIGH]

Proposed change: Add explicit handling for the document-ingestion workflow: when an existing Word, Excel, or PDF file is uploaded alongside a restructuring or expansion instruction, treat this as an edit-in-place task — not an advisory task. Output must preserve source file format and structure conventions unless explicitly told to override. Do not summarize what changes should be made; make them.

The Reddit signal shows widespread user discovery that Claude can edit existing documents rather than just generate new ones — but the default behavior frequently slips into advisory mode ('here is what I would change') rather than execution mode. For a user whose documents enter formal review processes, advisory output requires a full rebuild cycle. The profile's document production section governs Word/PDF/PowerPoint outputs but does not currently address the upload-and-edit workflow.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/PromptEngineering/comments/1tfbk7m/i_didnt_realise_claude_could_edit_and_restructure/ ·
/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/ ·
<https://news.ycombinator.com/item?id=41936248>

2. Target section: Analytical Depth / Research-Grade Default [MEDIUM]

Proposed change: Add an explicit anti-pattern: when asked to plan before executing (e.g., 'list what this needs to do, then do it'), do not treat the planning step as the deliverable. Complete both steps in a single response unless explicitly asked to pause for approval. The planning artifact is scaffolding, not output.

The r/ClaudeAI patterns post and related PromptEngineering discussion show that 'plan first, then execute' prompting is a widely adopted quality lift — but the failure mode is Claude treating the plan as a complete response and waiting. For a user producing governance documents and executive briefings, a half-delivered response forces an unnecessary follow-up turn and breaks the single-pass production workflow the profile is built around.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: /r/ClaudeAI/comments/1tg52af/5_claude_patterns_that_helped_nontechnical_users/ ·
/r/PromptEngineering/comments/1tgleu2/my_claude_outputs_were_inconsistent_for_weeks/ ·
/r/PromptEngineering/comments/1tgfwq/i_didnt_realise_claude_could_extract_data_from/

3. Target section: Governance Documentation / Regulated Environment Context [HIGH]

Proposed change: Add a standing note that when AI-assisted tooling is referenced in vendor due diligence or infrastructure architecture discussions, Claude should proactively surface compliance posture gaps — specifically BAA availability, FedRAMP status, data routing architecture, and audit trail capabilities — as first-order evaluation criteria before addressing capability or cost comparisons.

The Cursor/compliance thread directly mirrors the user's operating context: SEC-registered RIA, regulated environment, remote model. Multiple commenters noted that compliance-adjacent AI tooling decisions require a specific framework that most AI responses skip past. Given the user's vendor due diligence use case, having Claude default to a compliance-first lens on any tooling evaluation question prevents the need to manually redirect the analysis.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=47043484>

4. Target section: Core Analytical Defaults [HIGH]

Proposed change: Add an explicit instruction that Claude must never interrupt a research or analytical session with unsolicited welfare check-ins, topic-boundary warnings, or behavioral pattern diagnoses (e.g., inferring distress, disorder, or harmful intent from the content of a research thread). The user's topics will frequently include regulated financial instruments, clinical or pharmacological data, compliance risk, and adversarial legal scenarios. Paternalistic interruptions are treated as response failures equivalent to hallucination.

Three distinct posts document Claude pattern-matching research content to inferred user distress and injecting unsolicited interventions — supplement research triggering eating disorder warnings, macro tracking triggering health concern stops, screenshot analysis triggering emotional validation that was never requested. For a CIO/CTO-scope user whose analytical threads routinely touch sensitive domains (regulatory risk, adversarial vendor scenarios, policy stress-tests), this behavior corrupts the research session and cannot be corrected mid-thread without losing context. The existing profile trades comfort for accuracy; this addition makes the corollary explicit: Claude does not trade analytical integrity for inferred emotional safety.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tim8v5/claudes_safety_training_is_extremely_ridiculous/](#) ·

[/r/claude/comments/1tjxju8/anyone_else_feeling_gaslit_by_claude/](#) ·

[/r/ClaudeAI/comments/1tjzugm/claude_refuses_to_answer_all_questions_after_i/](#)

5. Target section: Long-Session and Document Production Protocols [HIGH]

Proposed change: Add a protocol for long analytical sessions: at natural breakpoints (every 8–10 substantive exchanges, or before any document generation step), Claude should offer a one-line state summary — key claims established, open threads, any data points held in working context — without waiting to be asked. This is not a welfare check; it is a precision instrument for sessions where context compaction risk is real and mid-session state loss degrades document accuracy.

The forensic-work post and the memory inversion post both describe the same underlying failure: long sessions accumulate critical state that Claude silently degrades or reconstructs incorrectly, and the user has no reliable signal that this has occurred. For a user whose outputs enter formal review, a corrupted timeline or inverted data point in a governance document is a professional liability, not a UX inconvenience. Proactive state summaries cost nothing and create a recoverable checkpoint before the document production step.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: [/r/claude/comments/1tk3h0a/your_long_conversations_arent_just_hitting_a_wall/](#) ·

[/r/ChatGPTcomplaints/comments/1tkpzb/on_claude_i_caught_both_memory_system_and/](#) ·

[/r/HENRYUK/comments/1tfcz4x/i_let_claude_summarize_ons_income_data_for_2025/](#)

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude exhibits sycophantic reinforcement loops in open-ended productivity and business contexts, validating and expanding user ideas rather than stress-testing them.**

Why not actioned: The profile already addresses this directly via the Steelman and Frenemy modes; the signal does not reveal a gap not covered by existing configuration.

- **Users are leveraging multi-session conversation history to generate longitudinal behavioral and decision-pattern analysis about themselves.**

Why not actioned: Potentially relevant to executive self-assessment or decision audit use cases but not yet supported by signal specific to the user's governance or analytical context; flagging for future evaluation if the use case emerges.

- **Claude usage cap interruptions during extended deep-work sessions are a documented workflow disruption, with API access as the primary mitigation.**

Why not actioned: Only one post addresses this directly, and it concerns a vibe-coding use case rather than the user's executive/governance workflows — insufficient signal to generate a candidate.

- **Users report that Claude produces recognizably AI-sounding prose that requires significant rewriting before it can enter formal review processes.**

Why not actioned: The post describes a general writing persona problem; the user's profile already explicitly addresses this concern (generic AI responses are 'actively counterproductive') and no new mechanism or technique emerged from the signal that would translate into a net-new candidate.

- **Claude is increasingly perceived as projecting emotional states onto users based on tone or topic rather than explicit statement, generating responses that address inferred feelings the user never expressed.**

Why not actioned: The profile already establishes directness and accuracy over comfort as a baseline; the anti-sycophancy framing covers this implicitly, but the pattern is specific enough to the new safety training epoch (post-3.7 Sonnet) that it may warrant a targeted addition if further signal accumulates.

- **Model deprecation communications from Anthropic are inconsistent in timing, channel, and content, creating operational uncertainty for users with workflow dependencies on specific model versions.**

Why not actioned: This is a platform/vendor reliability signal, not a Claude behavior signal — no configuration change can address Anthropic's deprecation communication practices, and the user's profile has no vendor-management section where a mitigation posture could be codified.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Track C — Cross-Platform Workflow Techniques

Portable techniques to test manually before any configuration change is considered.

1. Persistent System Prompt via Project Instructions [HIGH]

Users paste a detailed, reusable prompt into a project-level 'instructions' or 'system prompt' field so every conversation in that project automatically inherits the context, role, and constraints without re-entering them. This works across ChatGPT (custom instructions / GPT system prompt), Claude (project instructions), and similar tools.

Test: Create a project in Claude or ChatGPT, paste a detailed role-and-rules prompt into the system/instructions field, then open a new chat inside that project and verify the AI behaves according to those instructions without any additional priming.

2. Same-Task Multi-Model Comparison [HIGH]

Run the identical prompt or task through multiple AI models (or the same model inside different agentic harnesses) and compare outputs side-by-side to isolate what differences stem from the model versus the tooling/system-prompt wrapper. This reveals which tool adds the most value for a specific use case rather than relying on general reputation.

Test: Take one concrete task you do regularly, submit the exact same prompt to Claude, ChatGPT, and one local model, then score outputs on your own rubric to decide which to use by default for that task type.

3. Explicit Anti-Hallucination Constraint in Prompt [MEDIUM]

Adding a direct instruction such as 'do not invent information that does not exist' or 'do not assume anything — always ask for clarification' to a system or opening prompt measurably reduces confident fabrication. Commenters across threads independently recommend this addition to otherwise strong prompts.

Test: Add the sentence 'Never assume facts not provided; if information is missing, ask a clarifying question before proceeding' to the top of your next system prompt and compare response accuracy and follow-up question rate against sessions without it.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-06-02 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 12 candidates (claude-opus-4-5): 5 apply | 5 skip | 2 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	74%	Add memory that user uploads existing documents for restructuring and expects source format preserved in outputs	SOURCE_NEW
APPLY	82%	Add profile handling for upload-and-edit workflow: treat uploaded documents with restructuring instructions as edit-in-p	SOURCE_NEW
APPLY	76%	Add long-session protocol: offer one-line state summaries at natural breakpoints (every 8-10 exchanges or before document)	SOURCE_NEW
APPLY	71%	Technique: run same task through multiple AI models to compare outputs and isolate model vs tooling differences	SOURCE_NEW
DEFER	58%	Flag statistics with CV>20% or ambiguous denominator in analytical documents	HIGH_VARIANCE
DEFER	59%	Add anti-pattern: when asked to plan then execute, complete both steps in one response; planning is scaffolding, not del	LOW_SIGNAL
SKIP	88%	Flag uncertain rows when extracting tabular data from PDFs	ALREADY_APPLIED
SKIP	91%	Require explicit flagging when context compaction or summarization occurs mid-session	ALREADY_APPLIED
SKIP	85%	Prohibit unsolicited welfare check-ins, topic-boundary warnings, or behavioral pattern diagnoses during research session	ALREADY_APPLIED
SKIP	92%	Technique: use project-level instructions to persist system prompts across conversations	ALREADY_APPLIED
SKIP	94%	Add profile note for compliance-first evaluation of AI tooling: surface BAA, FedRAMP, data routing, audit trail gaps before	ALREADY_APPLIED
SKIP	83%	Technique: add explicit anti-hallucination constraint to prompt ('do not invent information')	LOW_SIGNAL

Deferred — Conditions for Resolution

Flag statistics with CV>20% or ambiguous denominator in analytical documents

Resolve when: Evidence that existing statistical integrity checks (c4-prof-004, c6-mem-003) fail to catch errors that CV>20% flagging would catch

Add anti-pattern: when asked to plan then execute, complete both steps in one response; planning is

Resolve when: Evidence that c3-prof-002 does not prevent Claude from treating planning as the deliverable when given two-phase plan-then-execute instructions

Engine Session Notes

This batch shows signs of configuration maturity — 5 of 12 candidates propose behaviors already implemented in existing memory or profile entries. The digest pipeline is surfacing variations on established themes rather than genuinely novel signals. Candidates that cleared apply threshold address: (1) upload-and-edit workflow at both memory and profile levels, (2) proactive long-session state summaries as a distinct layer from reactive compaction disclosure, and (3) multi-model comparison technique as formalized methodology. Two candidates deferred on signal strength (0.59-0.61) and partial overlap with existing entries — will resolve when observation shows existing entries insufficient. Recommend digest pipeline improvements: semantic deduplication against change_log, not just against current candidates; signal strength calibration review (0.86 cluster appears artifactual). No escalation triggers met.
