

CLAIRE

Claude Configuration Signal Report — 2026-06-07

Signal Summary

Run date: 2026-06-07

Posts scanned: 132

Signal Breakdown

Category	Count
behavior_complaint	22
workflow_gap	16
feature_praise	49
competitor_gap	1
cross_platform	0
noise (dropped)	44

Developer Signal Disposition

80 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Track A — Claude Native Signal

Memory Edit Candidates

No memory edit candidates this cycle. Signal below evidence threshold or no patterns identified.

Profile Diff Candidates

1. Target section: Vendor Due Diligence / Governance Documentation [HIGH]

Proposed change: Add a repeatable due diligence sub-framework for AI tooling in regulated environments. When evaluating AI vendors, Claude should default to surfacing: BAA availability, FedRAMP or equivalent certification status, data residency options, API vs. product-tier data handling distinctions, audit trail capabilities, and relevant compliance gap analysis (SOC2, HIPAA, SEC). Output should be structured for handoff to a compliance team without rewrite.

Two posts (43933054, 47043484) show a recurring pattern of regulated-industry practitioners struggling to evaluate AI tools against compliance requirements — exactly the vendor due diligence scope this configuration already covers. The signal is specific enough to warrant a named sub-framework rather than relying on general due diligence posture. Given the user's SEC-registered RIA context, this is a directly applicable gap.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=43933054> · <https://news.ycombinator.com/item?id=47043484>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude underperforms on long-horizon autonomous task completion where sequential decision-making and state tracking are required, distinct from single-turn analytical depth.**

Why not actioned: Only two posts touch this (43461944, 38382067) and neither maps to a use case in the user's configuration — the profile covers document production and analysis, not agentic task loops.

- **Claude capacity throttling and service availability issues are a recurring operational complaint, with multiple posts describing usage limits, rate constraints, and outages interrupting professional workflows.**

Why not actioned: Infrastructure reliability is a platform constraint outside the scope of configuration or prompt improvement — no memory edit or profile change can address availability or rate-limiting behavior.

- **Users express frustration with over-refusal and safety guardrails producing unhelpful responses, with some noting improvement in newer model versions but skepticism persisting.**

Why not actioned: The existing profile already explicitly addresses this concern — the configuration trades comfort for accuracy and instructs against hedge-everything responses; no new signal reaches the evidence threshold for a net-new actionable candidate distinct from what is already configured.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-06-07 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 1 candidates (claude-opus-4-5): 0 apply | 1 skip | 0 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
SKIP	78%	Add a repeatable due diligence sub-framework for AI vendors in regulated environments, surfacing BAA availability, FedRA	ALREADY_APPLIED, SOURCE_NEW

Engine Session Notes

Single-candidate batch. The candidate proposes vendor due diligence enhancements that substantially overlap with changes applied in cycles 5 and 6 (c5-prof-004, c6-prof-001, 2026-06-02 skill_install). Source post 47043484 has driven multiple prior candidates. Decision to skip reflects existing coverage rather than signal invalidity. The empty eval_history means hypothesis validation for prior decisions in this domain is pending - if future evals show vendor due diligence gaps, this skip decision should be revisited. Pipeline note: the digest layer may benefit from stronger deduplication against recently-applied changes, as this candidate's core proposal is already in the configuration.
