

CLAIRE

Claude Configuration Signal Report — 2026-06-14 — Cycle 8

Signal Summary

Run date: 2026-06-14

Posts scanned: 133

Signal Breakdown

Category	Count
behavior_complaint	19
workflow_gap	11
feature_praise	48
competitor_gap	2
cross_platform	1
noise (dropped)	52

Developer Signal Disposition

77 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Official Signal — Capability Deltas

Official-signal lane disabled in config this run.

Track A — Claude Native Signal

Memory Edit Candidates

No memory edit candidates this cycle. Signal below evidence threshold or no patterns identified.

Profile Diff Candidates

1. Target section: Vendor Due Diligence / Governance Documentation [HIGH]

Proposed change: Add a standing analytical frame for AI tooling procurement in regulated environments. When evaluating AI-adjacent vendor decisions (coding assistants, LLM APIs, agentic

platforms), default to a compliance-first risk taxonomy: data residency and routing, BAA/contractual coverage, certification status (SOC2, FedRAMP, HIPAA), audit trail adequacy, and shadow-IT exposure vectors. Surface gaps against the user's SEC-RIA regulatory posture without being asked.

The Cursor compliance post directly surfaces a vendor evaluation pattern — regulated-industry teams attempting to adopt AI tooling without adequate compliance scaffolding. The user operates an SEC-registered RIA in a regulated environment where this exact dilemma (AI productivity tools vs. compliance gatekeeping) is live. A standing frame prevents Claude from treating these as generic vendor questions rather than regulated-environment procurement decisions.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=47043484>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Community frustration with LLMs producing detectable, generic-voiced prose that requires significant rewriting before use — particularly in structured long-form outputs.**

Why not actioned: Supported by only one post with meaningful signal (hn_38382067); the other two posts do not corroborate a prose-quality or voice-drift pattern relevant to this user's configuration.

- **Skepticism about agentic AI reliability for sustained, multi-step autonomous task completion — with community consensus that current models fail on tasks requiring persistent state and sequential decision-making.**

Why not actioned: Supported by only one post (hn_43461944) and pertains to agentic/autonomous operation contexts not reflected in this user's established use cases; insufficient signal to action.

- **Anthropic service reliability complaints — outages, throttling, and capacity constraints are recurring friction points for heavy professional users.**

Why not actioned: Infrastructure availability is a platform constraint outside the scope of configuration or memory edits; no actionable profile change addresses it.

- **Community skepticism about Anthropic marketing claims versus actual model capability — particularly around safety-capability framing and benchmark comparisons.**

Why not actioned: Signal reflects general community sentiment about Anthropic as a company, not a behavioral pattern in Claude's outputs addressable through configuration.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-06-14 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 1 candidates (claude-opus-4-5): 0 apply | 1 skip | 0 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
SKIP	88%	Add a standing compliance-first analytical frame for AI vendor procurement decisions in regulated environments, covering	ALREADY_APPLIED, SOURCE_NEW

Engine Session Notes

Single-candidate batch representing pipeline noise. The digest generated a candidate from HN/47043484 that has already been processed and acted upon in Cycle 5 (c5-prof-004) using the same source. This suggests the semantic deduplication layer is not catching cross-cycle duplicates when phrasing differs. The fingerprint mechanism appears to be producing new fingerprints for semantically identical proposals. Recommend: add change_log source-URL matching to the assembler's suppression logic to catch candidates derived from already-processed source posts.