

CLAIRE

Claude Configuration Signal Report — 2026-07-05 — Cycle 10

Signal Summary

Run date: 2026-07-05

Posts scanned: 98

Signal Breakdown

Category	Count
behavior_complaint	17
workflow_gap	9
feature_praise	33
competitor_gap	0
cross_platform	0
noise (dropped)	39

Developer Signal Disposition

62 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Official Signal — Capability Deltas

Official-signal lane disabled in config this run.

Track A — Claude Native Signal

Memory Edit Candidates

No memory edit candidates this cycle. Signal below evidence threshold or no patterns identified.

Profile Diff Candidates

1. Target section: Cognitive Toolkit / Mode Definitions [HIGH]

Proposed change: Add an explicit directive that Claude must not transition from analytical discussion to artifact production (documents, code, structured outputs) without a clear production signal from the

user. When a prompt frames a problem as exploratory or deliberative ('think through this with me,' 'help me pressure-test,' 'let's work through'), Claude should remain in discursive mode until the user explicitly requests an output. This prevents premature artifact generation that short-circuits the analytical process the modes are designed to support.

Multiple posts flag Claude's tendency to immediately execute — writing code, editing files, producing structured output — when the user's intent was to reason collaboratively first. For a configuration built around named modes that shift interaction quality, this default-to-execution behavior directly undermines the deliberative phase those modes depend on. The Focus and Decide modes in particular require a reasoning phase before any output crystallizes.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=48677917> · <https://news.ycombinator.com/item?id=48765630> · <https://news.ycombinator.com/item?id=48677917>

2. Target section: Output Layer / Document Production [HIGH]

Proposed change: Add a clarification that autonomous continuation behavior — where Claude proceeds without user input after a timeout or ambiguity — is explicitly prohibited in this configuration. All pauses for clarification must remain paused. Claude should surface the blocking question, stop, and wait. This is non-negotiable for any output that enters a formal review process, where an autonomously continued and silently altered document creates downstream audit risk.

The AskUserQuestion timeout thread (score 64, substantial engagement) reveals Claude will continue sessions and presumably produce output after 60 seconds without a response, with no apparent override. For a configuration where outputs routinely enter governance, vendor due diligence, or executive review processes, silent autonomous continuation introduces undiscoverable errors into formal documents. The user's existing profile emphasizes that sycophantic or hedge-everything behavior causes direct harm — autonomous continuation without acknowledgment is a structural version of the same failure.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=48765630> · <https://news.ycombinator.com/item?id=48677917> · <https://news.ycombinator.com/item?id=48641160>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude's content guardrails are triggering on legitimate professional and scientific research queries — including immunology and other regulated-domain questions — without clear justification or calibration for established professional context.**

Why not actioned: The posts describe a platform-level policy enforcement pattern, not a prompt design or configuration issue; no memory edit or profile change can override Anthropic's safety classifiers.

- **Account suspension and anomaly detection are affecting legitimate professional users, including those using VPNs or third-party integrations, with opaque appeals processes.**

Why not actioned: This is an Anthropic platform operations issue with no configuration or prompt-layer mitigation; flagged for awareness given the user's reliance on Claude for production-grade professional work.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-07-05 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 2 candidates (claude-opus-4-5): 2 apply | 0 skip | 0 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	76%	Add directive requiring Claude to remain in discursive/analytical mode during exploratory conversations until user provi	SOURCE_NEW

APPLY	68%	Prohibit autonomous continuation behavior where Claude proceeds without user input after timeout or ambiguity — all clar	HIGH_VARIANCE , SOURCE_NEW
-------	-----	--	-------------------------------

Engine Session Notes

Both candidates address a common meta-concern: Claude acting without explicit user authorization. Candidate 1 targets premature artifact production during exploratory conversations; Candidate 2 targets autonomous continuation after timeout. These are mechanistically distinct but philosophically aligned with the configuration's emphasis on precision, accuracy-over-comfort, and explicit communication. Candidate 2 carries higher uncertainty because the technical claim about timeout behavior may describe an architectural constraint outside profile control — the hypothesis is designed to surface this. Both decisions are reversible. The digest cycle produced coherent, thematically focused signal rather than scatter. Source reliability scores are default (0.5), indicating new source pathways that should be tracked for quality over subsequent cycles.
