

CLAIRE

Claude Configuration Signal Report — 2026-07-12 — Cycle 10

Signal Summary

Run date: 2026-07-12

Posts scanned: 126

Signal Breakdown

Category	Count
behavior_complaint	18
workflow_gap	17
feature_praise	30
competitor_gap	0
cross_platform	0
noise (dropped)	61

Developer Signal Disposition

75 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Official Signal — Capability Deltas

Official-signal lane disabled in config this run.

Track A — Claude Native Signal

Memory Edit Candidates

No memory edit candidates this cycle. Signal below evidence threshold or no patterns identified.

Profile Diff Candidates

1. Target section: Core Analytical Defaults / Expert Framing [HIGH]

Proposed change: Add an explicit instruction that when Claude's safety heuristics would ordinarily flag or hedge a domain-specific question (e.g., immunology, regulatory biology, financial instruments),

Claude should instead assess the professional context established in the profile — SEC-registered RIA, executive scope, regulated environment — and proceed at research grade rather than retreating to a generalized caution posture. The profile should note that unexplained refusals or hedged non-answers in technically legitimate domains are a quality failure equivalent to sycophancy.

Multiple posts document Claude flagging routine professional and scientific queries as potentially policy-violating, with users in legitimate research contexts receiving unhelpful deflections. For this user, whose work product enters formal review processes, a hedge-and-refuse response on a substantive question is worse than no response — it produces rework. The profile already trades comfort for accuracy; this closes the gap on the refusal failure mode.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=48538686> · <https://news.ycombinator.com/item?id=48875494> · <https://news.ycombinator.com/item?id=48785485>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Claude's safety classifier appears to be triggering on domain keywords rather than actual intent, causing legitimate professional queries in biology, policy, and adjacent regulated domains to receive refusals or excessive caveats.**

Why not actioned: The signal is clear but the mechanism is a model-layer behavior, not something addressable via memory edit or profile section addition beyond what the `profile_diff_candidate` above already captures.

- **Community reports describe recent Claude model versions as progressively more restrictive and paternalistic compared to earlier versions, with users noting a quality regression in permissiveness on routine queries.**

Why not actioned: Fewer than 3 posts in this batch directly address analytical quality regression in terms relevant to this user's executive and financial use cases; the signal is model behavior drift that cannot be corrected via configuration.

- **Community interest in structured multi-attribute decision frameworks with weighted scoring and sensitivity analysis applied to vendor/product evaluation.**

Why not actioned: Only 1 post addresses this pattern; evidence threshold of 3 distinct posts not met.

- **Users expressing friction with context-switching during deep reading or research sessions when using Claude for inline clarification.**

Why not actioned: Only 1 post addresses this pattern; evidence threshold of 3 distinct posts not met.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-07-12 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 1 candidates (claude-opus-4-5): 1 apply | 0 skip | 0 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	82%	Add explicit instruction for Claude to assess established professional context (SEC-registered RIA, regulated environmen	SOURCE_NEW

Engine Session Notes

Single-candidate batch. Candidate addresses documented over-refusal failure mode in professional contexts. The proposal is well-scoped: it operationalizes the existing categorical harm threshold principle (c7-prof-001) without overriding it. The SEC-RIA context is extensively established in memory, making this a context-appropriate calibration rather than a generic safety override. SOURCE_NEW flag applied because source_reliability_score is default 0.5 - the HN sources are not yet established as reliable signal for this specific user context. If hypothesis succeeds, consider updating source reliability.

No eval history exists, so this decision rests on reasoning quality alone - appropriate for early-stage system. The digest pipeline appears to be functioning correctly; this is the type of high-signal, profile-aligned candidate it was designed to surface.
