

CLAIRE

Claude Configuration Signal Report — 2026-07-19 — Cycle 10

Signal Summary

Run date: 2026-07-19

Posts scanned: 164

Signal Breakdown

Category	Count
behavior_complaint	30
workflow_gap	15
feature_praise	45
competitor_gap	0
cross_platform	0
noise (dropped)	74

Developer Signal Disposition

97 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Official Signal — Capability Deltas

Official-signal lane disabled in config this run.

Track A — Claude Native Signal

Memory Edit Candidates

No memory edit candidates this cycle. Signal below evidence threshold or no patterns identified.

Profile Diff Candidates

1. Target section: Core Analytical Defaults [HIGH]

Proposed change: Add an explicit instruction that Claude should never preemptively shift into execution or output-production mode when the user's framing signals exploratory or deliberative intent. Keywords

like 'think through,' 'help me reason about,' 'let's consider,' or 'what's your read on' should trigger sustained analytical dialogue, not deliverable production. Claude should surface its reasoning and wait for explicit direction before drafting, structuring, or producing a document artifact.

Multiple posts report Claude defaulting to immediate action or output generation when users intended collaborative reasoning. For a profile where analytical depth is the default operating mode, premature execution collapses the deliberative phase that produces the highest-quality inputs to decisions. This is especially harmful in governance, architecture, and business case work where the reasoning process is as important as the output.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=48677917> · <https://news.ycombinator.com/item?id=48875494> · <https://news.ycombinator.com/item?id=48918956>

2. Target section: Accuracy and Honesty Standards [HIGH]

Proposed change: Add a note that Claude should proactively flag when a topic area — including regulated or sensitive domains (life sciences, certain geopolitical actors, financial instrument specifics) — may be subject to classifier over-triggering or model-level conservatism that could degrade response quality. If Claude detects it is approaching a topic where it has historically hedged inappropriately or where its training may produce degraded signal, it should name that limitation explicitly rather than silently softening the analysis.

Community signal across multiple posts documents Claude refusing or degrading responses in legitimate professional contexts — biology research, life sciences, regulatory topics — without signaling why. For a profile operating in regulated environments where analytical gaps have direct decision consequences, silent degradation is more dangerous than an explicit limitation disclosure. The user's configuration already trades comfort for accuracy; this extends that principle to meta-level honesty about Claude's own reliability.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=48538686> · <https://news.ycombinator.com/item?id=48954522> · <https://news.ycombinator.com/item?id=48875494>

3. Target section: Output and Tone Defaults [HIGH]

Proposed change: Add an explicit guard against unsolicited moral commentary, safety caveats, or value-framing appended to analytical responses. If the user's query is professional, contextually grounded, and does not request ethical evaluation, Claude should not volunteer it. Judgment about whether a topic warrants ethical framing should be based on actual content risk, not surface-level topic proximity to sensitive categories.

Multiple posts describe Claude becoming increasingly judgmental and prone to unprompted lecturing even on routine professional or analytical queries. For a profile where sycophantic or hedge-everything responses are explicitly counterproductive, moralistic padding is the inverse failure mode — it signals low analytical confidence and erodes trust in outputs that enter formal review processes. The configuration already requires direct and honest responses; this closes the gap where that principle is violated by tone rather than content.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=48918956> · <https://news.ycombinator.com/item?id=48875494> · <https://news.ycombinator.com/item?id=48538686>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Community reporting a pattern of account suspensions triggered by anomaly detection heuristics (VPN usage, shared payment methods, geographic signals) rather than actual policy violations, with opaque support responses.**

Why not actioned: This is a platform trust and account management issue, not a configuration or behavior pattern that can be addressed through profile or memory edits.

- **Claude Code usage limits are being reduced and access to frontier models (Opus/Fable-tier) is being restricted to higher pricing tiers, with community frustration about rate limit unpredictability.**

Why not actioned: Pricing and access tier decisions are Anthropic product policy, not addressable through configuration changes; worth monitoring if it affects the user's operational model access.

- **Community discussion around Claude Code as a terminal-native agentic workflow layer, distinct from web chat, with sysadmin-adjacent power users treating it as primary interface**

Why not actioned: Only one post in the input; evidence threshold requires 3 or more distinct posts to generate a candidate, and the post's sysadmin/non-IT framing does not map to the user's CIO/CTO executive or trading use cases

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

No profile addition candidates this cycle.

Eval Loop Status

Applied Changes Pending Eval

Review `change_log.json` for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-07-19 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 4 candidates (claude-opus-4-5): 3 apply | 1 skip | 0 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
APPLY	82%	Add instruction to detect deliberative intent from keywords and stay in analytical mode rather than shifting to delivera	SOURCE_NEW
APPLY	78%	Add instruction to proactively flag when approaching topic areas where classifier over-triggering or model conservatism	SOURCE_NEW
APPLY	80%	Add explicit guard against unsolicited moral commentary, safety caveats, or value-framing appended to analytical respons	SOURCE_NEW
SKIP	71%	User technique: define hard constraints and weighted priorities upfront before recommendations, AI scores options agains	LOW_SIGNAL, SOURCE_NEW

Engine Session Notes

Batch shows coherent signal around unwanted Claude insertions - premature execution, silent degradation, and moral padding are three variants of 'Claude adding content user didn't request.' Three HIGH-confidence profile candidates with 0.86 signal strength all address this theme from different angles and merit application. One MEDIUM-confidence technique candidate with 0.57 signal strength targets user behavior rather than profile configuration and correctly routes to skip - pipeline may benefit from earlier filtering by change_target to avoid routing user-behavior techniques to profile evolution decisions. No dependencies among candidates. All candidates have default source reliability (0.5) and zero prior appearances, indicating fresh batch from new signal source.
