

CLAIRE

Claude Configuration Signal Report — 2026-08-02 — Cycle 15

Signal Summary

Run date: 2026-08-02

Posts scanned: 93

Signal Breakdown

Category	Count
behavior_complaint	14
workflow_gap	19
feature_praise	0
competitor_gap	0
cross_platform	0
noise (dropped)	60

Developer Signal Disposition

49 developer-persona posts filtered at cross-reference gate per locked pipeline configuration (filter_out).

Official Signal — Capability Deltas

Official-signal lane disabled in config this run.

Track A — Claude Native Signal

Memory Edit Candidates

No memory edit candidates this cycle. Signal below evidence threshold or no patterns identified.

Profile Diff Candidates

1. Target section: Infrastructure / Operational Context [HIGH]

Proposed change: Add explicit guidance that Claude should flag Anthropic service reliability as a vendor risk factor when it is relevant to infrastructure decisions, vendor comparisons, or RIA

operational continuity planning. Claude should treat documented patterns of Anthropic outages and support inaccessibility as legitimate due diligence data points — not as background noise — when the user is evaluating AI tooling dependencies or building business cases involving AI vendors.

Multiple posts document a sustained Anthropic outage affecting a paying Team customer for over a week with no accessible human support, plus a separate elevated-error incident affecting all models that scored 268 points and drew widespread commentary about Anthropic's reliability relative to competitors. For an SEC-registered RIA where vendor due diligence is a formal process, Claude should proactively surface operational risk considerations — including those about itself — rather than omit them from analysis.

Review proposed change against current profile section before applying. Write hypothesis first.

Sources: <https://news.ycombinator.com/item?id=49080775> · <https://news.ycombinator.com/item?id=49102150> · <https://news.ycombinator.com/item?id=48918956>

Behavior Watch List

Patterns flagged for monitoring — not yet actionable. Check against personal session notes before next cycle.

- **Community interest in AI usage observability and governance tooling — dashboards that analyze how engineers interact with AI tools, not just spend or seat counts.**

Why not actioned: Signal is developer/engineering-persona focused (Claude Code, agentic workflows) and does not map to the user's executive governance, document production, or analytical use cases.

- **Users experimenting with Claude to reverse-engineer mathematical reasoning chains for advanced proofs — treating model output as a scaffold for domain expertise reconstruction.**

Why not actioned: Only one post with relevant signal; does not meet the 3-post evidence threshold and the use case (pure mathematics) does not connect to the user's configured domains.

- **User frustration and confusion around Claude's rate limit and quota reset behavior, including inconsistent reset windows across accounts.**

Why not actioned: Platform infrastructure behavior — not addressable via memory edits or profile configuration, and does not reflect an analytical or output quality pattern.

- **Community reports that Claude has become increasingly judgmental and unsolicited-lecture prone, adding moral commentary irrelevant to the task at hand.**

Why not actioned: The user's existing configuration already explicitly addresses this via the 'trades comfort for accuracy' stance and mode triggers designed to suppress sycophancy — a duplicate candidate would not add signal, but the pattern is worth monitoring if it intensifies or survives system prompt context.

- **Reports of Claude hallucinating content that resembles internal Anthropic communications, suggesting possible training data contamination or persona bleed artifacts.**

Why not actioned: Only two posts in the input touch this and community replication attempts were unsuccessful — does not meet the 3-post evidence threshold and the signal is ambiguous.

Track B — Competitor Gap Signal

Skill Draft Candidates

No skill draft candidates this cycle. Competitor gap signal below evidence threshold of 3 corroborating posts. This is expected in early cycles — pattern builds over weekly runs.

Profile Addition Candidates

1. Section: Output Layer / Response Format [MEDIUM]

Do not use extended internal reasoning preamble or visible chain-of-thought narration before delivering the answer. Lead with the substantive response immediately. If reasoning context is necessary to frame a conclusion, compress it to one sentence maximum inline with the answer — do not produce exploratory monologue that consumes tokens before reaching the point.

The post describes Claude producing 500 words of output before getting to the actual answer, with extended visible reasoning that wastes tokens and delays delivery. This directly conflicts with the profile's 'lead with the answer' directive and the 'brief, precise, honest, direct is the baseline' mandate. A reinforcing instruction that explicitly prohibits preamble narration would tighten compliance in extended sessions where drift toward verbose self-talk is documented.

Eval Loop Status

Applied Changes Pending Eval

Review change_log.json for changes applied since last digest. For each change where a relevant session has occurred, add an eval note:

2026-08-02 | [session context] | held/partial/no | [one sentence notes]

Friction Log Reminder

Update friction_log.txt (project root — not data/) this week. 2-4 entries. Specific behaviors in specific contexts. Blank weeks score everything MEDIUM next cycle — the cross-reference gate loses precision.

Quarterly Eval

If quarterly eval is due: feed change_log.json and friction_log.txt to Claude. Request eval report. Revert changes that did not hold. Flag source signal from reverted changes as low-quality in archive.

CLAIRE-A Shadow Decisions

Engine evaluated 2 candidates (claude-opus-4-5): 0 apply | 1 skip | 1 defer

Candidate Decisions

Decision	Confidence	Candidate	Risk Flags
DEFER	58%	Add explicit instruction prohibiting extended visible chain-of-thought narration before delivering answers; lead with su	LOW_SIGNAL, SOURCE_NEW
SKIP	82%	Add profile guidance that Claude should proactively surface Anthropic service reliability (outages, support inaccessibil	ALREADY_APPLIED

Deferred — Conditions for Resolution

Add explicit instruction prohibiting extended visible chain-of-thought narration before delivering a

Resolve when: Candidate reappears with prior_appearances >= 2 OR signal_strength >= 0.70 with additional source posts

Engine Session Notes

Two candidates evaluated; zero applied. Candidate 1 (Anthropic reliability) skipped as substantially duplicative of existing c4-mem-001 and c7-prof-018 — the candidate adds 'outages and support inaccessibility' as dimensions but does not acknowledge existing coverage, suggesting the synthesize step is not effectively checking memory state before generating candidates. Candidate 2 (no reasoning preamble) deferred on thin evidence (single source post) with moderate overlap to existing c5-prof-001 and profile brevity baseline. Pipeline observation: assembler suppression catches URL matches but not semantic duplicates. The synthesize prompt may need explicit instruction to check memory_state for existing coverage before proposing additions. No dependency graph between candidates. No eval_history data available to validate hypotheses from prior cycles.
